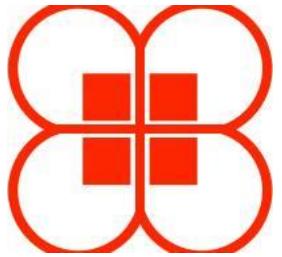


# Reproducible Performance Measurement and Analysis for High-Performance Computing Applications

HPDC Tutorial



July 20, 2025

Stephanie Brink, David Boehme, Ian Lumsden,  
Dewi Yokelson, Michela Taufer, and Olga Pearce



LLNL-CFPRES-2202755

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

 THE UNIVERSITY OF  
TENNESSEE  
KNOXVILLE

 Lawrence Livermore  
National Laboratory

# Tutorial Presenters

---



Stephanie Brink  
LLNL



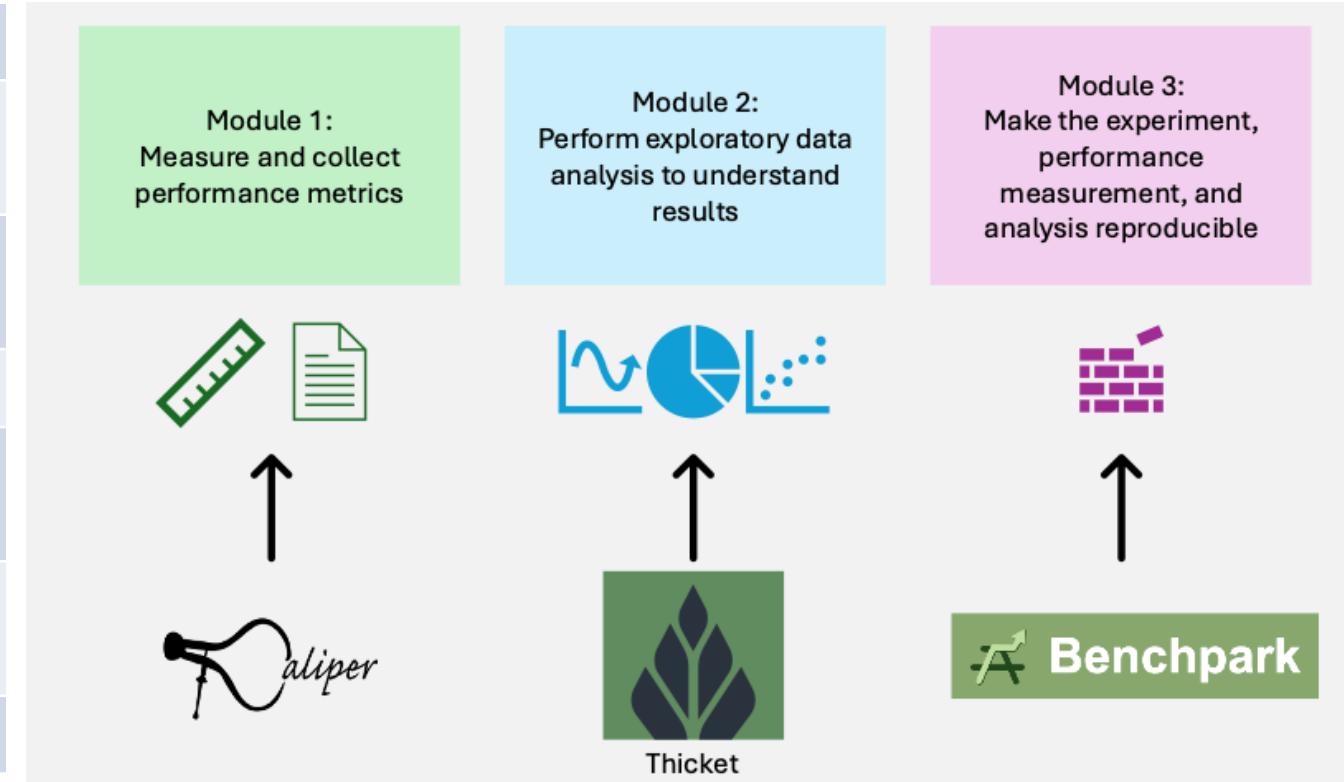
David Boehme  
LLNL



Ian Lumsden  
UTK

# Overall Tutorial Agenda (approximate)

Welcome and Introduction	15 min
Module 1: Performance measurement with Caliper	45 min
Module 2: Performance analysis with Thicket	30 min
Break	30 min
Module 2 (cont'd): Performance analysis with Thicket	20 min
Module 3: Reproducible experiments with Benchpark	60 min
Q&A and Wrapup	10 min



\*Hands on component to accompany each module

# Overall Tutorial Materials

---

Find these slides and associated scripts here:

- Caliper: <https://software.llnl.gov/Caliper/> under “Materials”
- Thicket: [https://thicket.readthedocs.io/en/latest/tutorial\\_materials.html](https://thicket.readthedocs.io/en/latest/tutorial_materials.html)
- Benchpark: <https://software.llnl.gov/benchpark/basic-tutorial.html>

We also have a channel on Spack slack.

You can join here:

**slack.spack.io**

**Join the #benchpark-support channel!**

You can continue to ask questions here after the tutorial has ended.

## Tutorial Instances: <http://bit.ly/4kGQDlc>

---

- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
  - Pre-installed Thicket, Caliper, and Benchpark and required dependencies
  - Caliper source code demos
  - Thicket Jupyter notebooks and datasets for performance analysis



When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
  - First initial followed by last name (e.g., jdoe)
- PW: hpctutorial25

---

# Performance Measurement with Caliper

David Boehme



# Caliper: A Performance Profiling Library

2025 HPC Tutorial

July 20, 2025

David Boehme  
Computer Scientist



# Caliper: A Performance Profiling Library

---

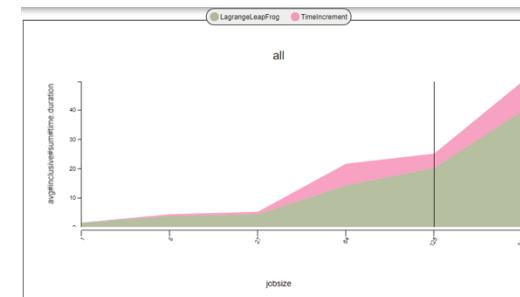
- Integrates a performance profiler into your program
  - Profiling is always available
  - Simplifies performance profiling for application end users
- Common instrumentation interface
  - Provides program context information for other tools
- Designed for HPC
  - Supports MPI, OpenMP, CUDA, HIP, Kokkos, RAJA

# Caliper Use Cases

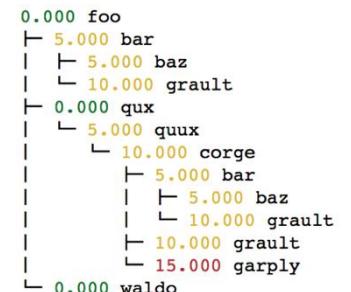
- Lightweight always-on profiling
  - Performance summary report for each run
- Performance debugging
- Performance introspection
- Comparison studies across runs
  - Performance regression testing
  - Configuration and scaling studies
- Automated workflows

Performance reports

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
main	0.000119	0.000119	0.000119	7.079120
mainloop	0.000067	0.000067	0.000067	3.985723
foo	0.000646	0.000646	0.000646	38.429506
init	0.000017	0.000017	0.000017	1.011303

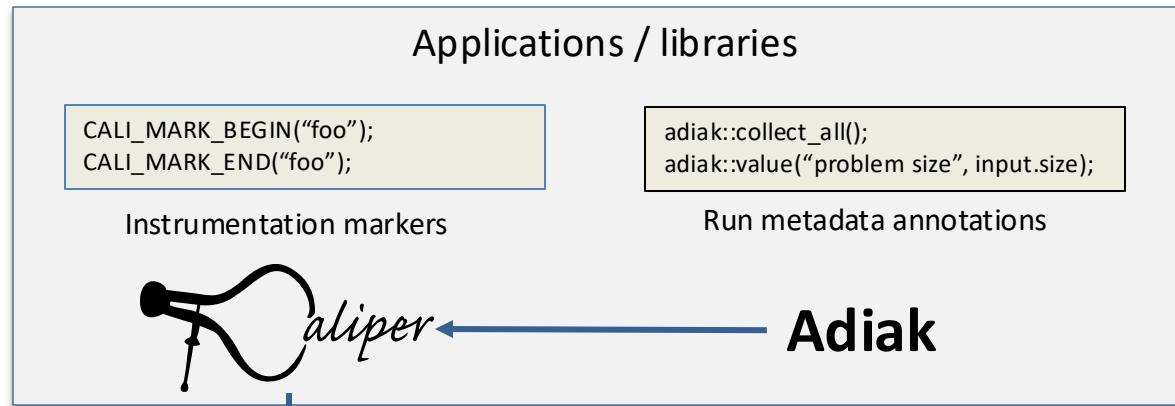


Comparing runs



Debugging

# The Caliper Performance Analysis Stack

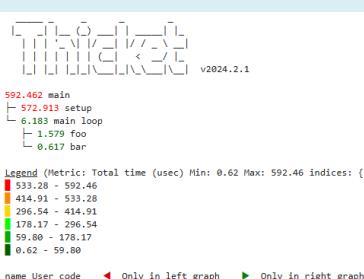


## Reports

Path	Time (E)	Time (I)	Time % (E)
main	0.000008	0.000379	1.953781
setup	0.000366	0.000366	84.122052
main loop	0.000003	0.000005	0.780592
foo	0.000001	0.000001	0.267251
bar	0.000001	0.000001	0.134545



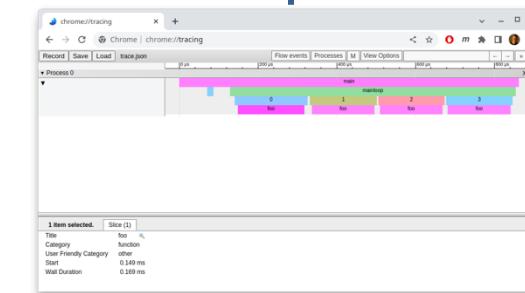
Thicket



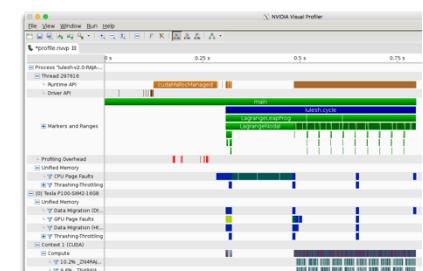
TreeScape

Python frameworks for analyzing large sets of runs

Event traces



Annotation bindings



3rd party/vendor tools

# Materials, Contact & Links

- Tutorial materials: <https://github.com/daboehm/caliper-tutorial>

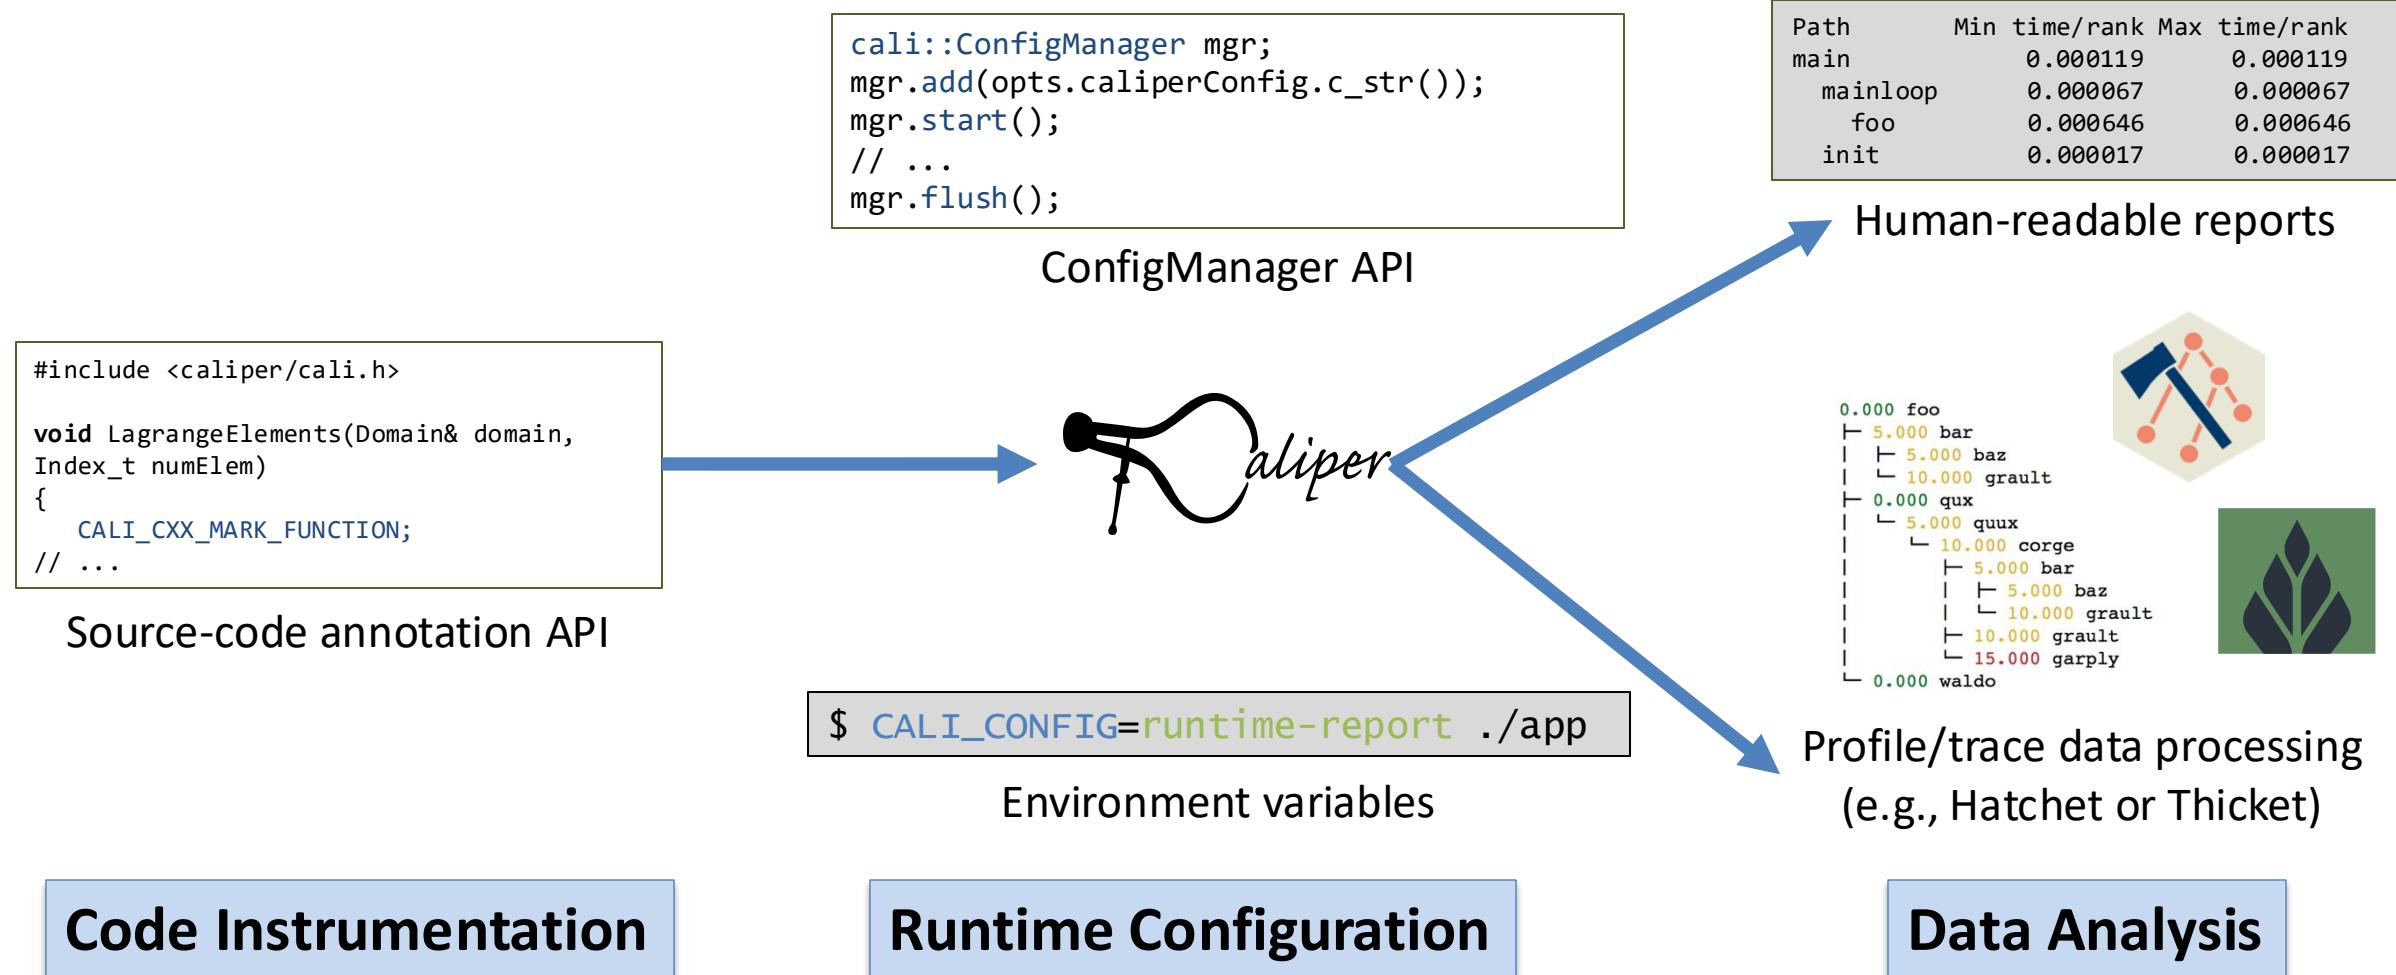
```
$ git clone --recursive https://github.com/daboehm/caliper-tutorial.git  
$ . setup-env.sh
```

- GitHub repository: <https://github.com/LLNL/Caliper>
- Documentation: <https://llnl.github.io/Caliper>
- GitHub Discussions: <https://github.com/LLNL/Caliper/discussions>
- Contact: David Boehme (boehme3@llnl.gov)

---

# Using Caliper

# Using Caliper: Workflow



Code Instrumentation

Runtime Configuration

Data Analysis

# Region Profiling: Marking Code Regions

C/C++

```
#include <caliper/cali.h>

void main() {
    CALI_MARK_BEGIN("init");
    do_init();
    CALI_MARK_END("init");

    CALI_MARK_PHASE_BEGIN("hydro");
    do_hydro();
    CALI_MARK_PHASE_END("hydro");
}
```

Fortran

```
USE caliper_mod

CALL cali_begin_region('init')
CALL do_init()
CALL cali_end_region('init')
```

- Use annotation macros (C/C++) or functions to mark and name code regions

# Region Profiling: Best Practices

- Be selective: Instrument high-level program subdivisions (kernels, phases, ...)
- Be clear: Choose meaningful names
- Start small: Add instrumentation incrementally

```
RAJA::ReduceSum<RAJA::omp_reduce, double> ompdot(0.0);

CALI_MARK_BEGIN("dotproduct");

RAJA::forall<RAJA::omp_parallel_for_exec>(RAJA::RangeSegment(0, N), [=] (int i) {
    ompdot += a[i] * b[i];
});
dot = ompdot.get();

CALI_MARK_END("dotproduct");
```

Caliper annotations give meaningful names to high-level program constructs

# Instrumentation Support for RAJA and Kokkos

```
RAJA::dynamic_forall<policy_list>(pol, range,
    RAJA::expt::Reduce<RAJA::operators::plus>(&sum),
    RAJA::Name("RAJA dynamic forall reduction"),
    [=] RAJA_HOST_DEVICE (int i, VAL_INT_SUM &_sum) {
        c[i] = a[i] + b[i];
        _sum += 1;
    });
}
```

- RAJA now supports Caliper via optional RAJA::Name() parameter to label a kernel
- Caliper can also be used as a Kokkos profiling tool

# Region Profiling: Printing a Runtime Report

```
$ cd Caliper/build  
$ make cxx-example  
$ CALI_CONFIG=runtime-report ./examples/apps/cxx-example
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
main	0.000119	0.000119	0.000119	7.079120
mainloop	0.000067	0.000067	0.000067	3.985723
foo	0.000646	0.000646	0.000646	38.429506
init	0.000017	0.000017	0.000017	1.011303

- Set the CALI\_CONFIG environment variable to access Caliper's built-in profiling configurations
- “runtime-report” measures, aggregates, and prints time in annotated code regions

# List of Caliper's Built-in Profiling Recipes

Config name	Description
runtime-report	Print a time profile for annotated regions
loop-report	Print summary and time-series information for loops
mpi-report	Print time spent in MPI functions
sample-report	Print time spent in regions using call-path sampling
event-trace	Record a trace of region enter/exit events in .cali format
hatchet-region-profile	Record a region time profile for processing with hatchet or cali-query
hatchet-sample-profile	Record a sampling profile for processing with hatchet or cali-query
spot	Record a time profile for Thicket or TreeScape

Use `cali-query --help=configs` to list all built-in configs and their options

# Built-In Profiling Recipes: Configuration String Syntax

*Config name* specifies the kind of performance measurement

*Parameters* enable additional features, metrics, or output options

```
$ CALI_CONFIG="runtime-report(mem.highwatermark,output=stdout)" ./examples/apps/cxx-example
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %	Allocated MB
main	0.000179	0.000179	0.000179	2.054637	0.000047
mainloop	0.000082	0.000082	0.000082	0.941230	0.000016
foo	0.000778	0.000778	0.000778	8.930211	0.000016
init	0.000020	0.000020	0.000020	0.229568	0.000000

- Most Caliper measurement recipes have optional parameters to enable additional features or configure output settings

# Profiling Options: MPI Function Profiling

```
$ CALI_CONFIG=runtime-report,profile.mpi ./lulesh2.0
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
MPI_Comm_dup	0.000034	0.003876	0.001999	0.10089
main	0.009013	0.010797	0.010173	0.51335
MPI_Reduce	0.000031	0.000049	0.000037	0.001886
lulesh.cycle	0.002031	0.002258	0.002085	0.105220
LagrangeLeapFrog	0.002158	0.002511	0.002227	0.112366
CalcTimeConstraintsForElms	0.015166	0.015443	0.015277	0.770922
CalcQForElms	0.058781	0.060196	0.059699	3.01254
CalcMonotonicQForElms	0.035331	0.041057	0.038496	1.942601
CommMonoQ	0.005280	0.006152	0.005544	0.279781
MPI_Wait	0.004182	0.084533	0.035324	1.78249
CommSend	0.006893	0.009062	0.008071	0.407298
MPI_Waitall	0.000986	0.001778	0.001343	0.067789
MPI_Isend	0.004564	0.005785	0.004930	0.248765
CommRecv	0.002265	0.002616	0.002341	0.118144
[...]				

The profile.mpi option measures time spent in MPI functions

# Profiling Options: CUDA Profiling

```
$ lrun -n 4 ./tea_leaf runtime-report,profile.cuda
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
timestep_loop	0.000175	0.000791	0.000345	0.002076
[...]				
total_solve	0.000105	0.000689	0.000252	0.001516
solve	0.583837	0.617376	0.594771	3.581811
dot_product	0.000936	0.001015	0.000969	0.005837
cudaMalloc	0.000060	0.000066	0.000063	0.000382
internal_halo_update	0.077627	0.079476	0.078697	0.473925
halo_update	0.158597	0.161853	0.160023	0.963685
halo_exchange	1.502106	1.572522	1.532860	9.231136
cudaMemcpy	11.840890	11.871018	11.860343	71.424929
cudaLaunchKernel	1.177454	1.230816	1.211668	7.296865
cudaMemcpy	0.470123	0.471485	0.470596	2.834008
cudaLaunchKernel	0.658269	0.682566	0.673030	4.053100
[...]				

The profile.cuda option measures time in CUDA runtime API calls

# Use cali-query to get Help for Profiling Recipes

```
$ cali-query --help runtime-report
```

## runtime-report

Print a time profile for annotated regions

### Options:

aggregate_across_ranks	Aggregate results across MPI ranks
calc.inclusive	Report inclusive instead of exclusive times
cuda.gputime	Report GPU time in CUDA activities
exclude_regions	Do not take snapshots for the given region names/patterns.
include_branches	Only take snapshots for branches with the given region names.
include_regions	Only take snapshots for the given region names/patterns.
io.bytes	Report I/O bytes written and read
io.bytes.read	Report I/O bytes read
io.bytes.written	Report I/O bytes written
io.read.bandwidth	Report I/O read bandwidth
io.write.bandwidth	Report I/O write bandwidth
level	Minimum region level that triggers snapshots
main_thread_only	Only include measurements from the main thread in results.
max_column_width	Maximum column width in the tree display
...	

# Sample Profiling

```
$ CALI_CONFIG=sample-report ./lulesh2.0
```

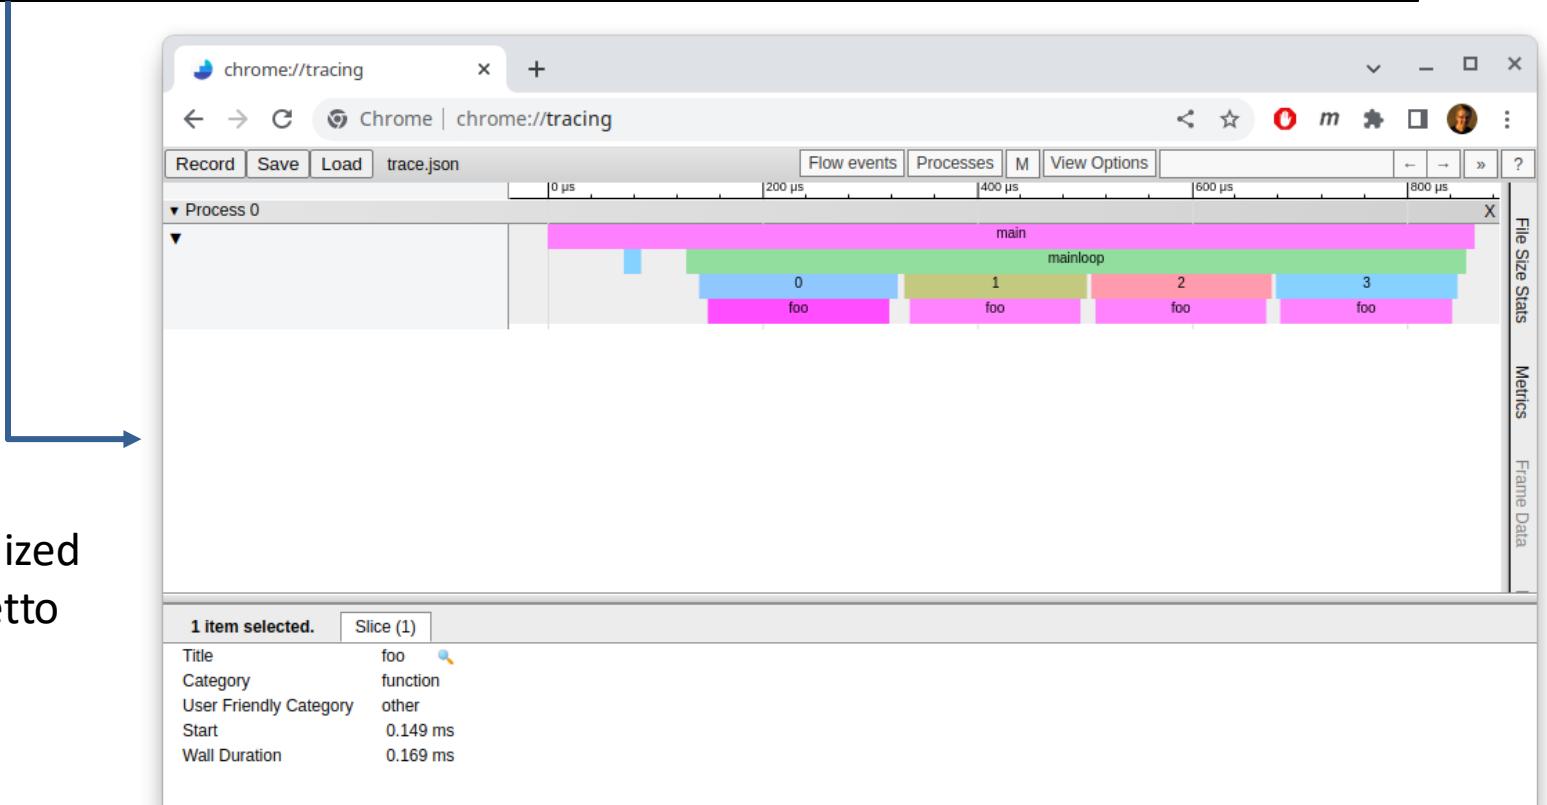


Path	Min time/rank	Max time/rank	Avg time/rank	Total time	Time %	Function
main						
-	0.005000	0.005000	0.005000	0.035000	0.059691	Domain::AllocateElemPersistent
-	0.005000	0.005000	0.005000	0.035000	0.059691	Domain::SetupThreadSupportStru
-	0.005000	0.005000	0.005000	0.005000	0.008527	sysmalloc
-	0.005000	0.005000	0.005000	0.005000	0.008527	Domain::BuildMesh(int, int, in
lulesh.cycle						
TimeIncrement						
-	0.075000	0.740000	0.355000	2.840000	4.843523	gomp_barrier_wait_end
-	0.005000	0.060000	0.027857	0.195000	0.332566	psm2_mq_ipeek2
-	0.005000	0.005000	0.005000	0.015000	0.025582	psm_no_lock
-	0.005000	0.060000	0.023571	0.165000	0.281402	psm_progress_wait
-	0.015000	0.030000	0.022143	0.155000	0.264347	mv2_shm_bcast
-	0.005000	0.025000	0.013750	0.055000	0.093801	amsh_poll
-	0.005000	0.010000	0.007500	0.030000	0.051164	psmi_poll_internal
...						

The *sample-report* recipe samples source functions or file+line locations.

# Event Tracing and Timeline Visualization

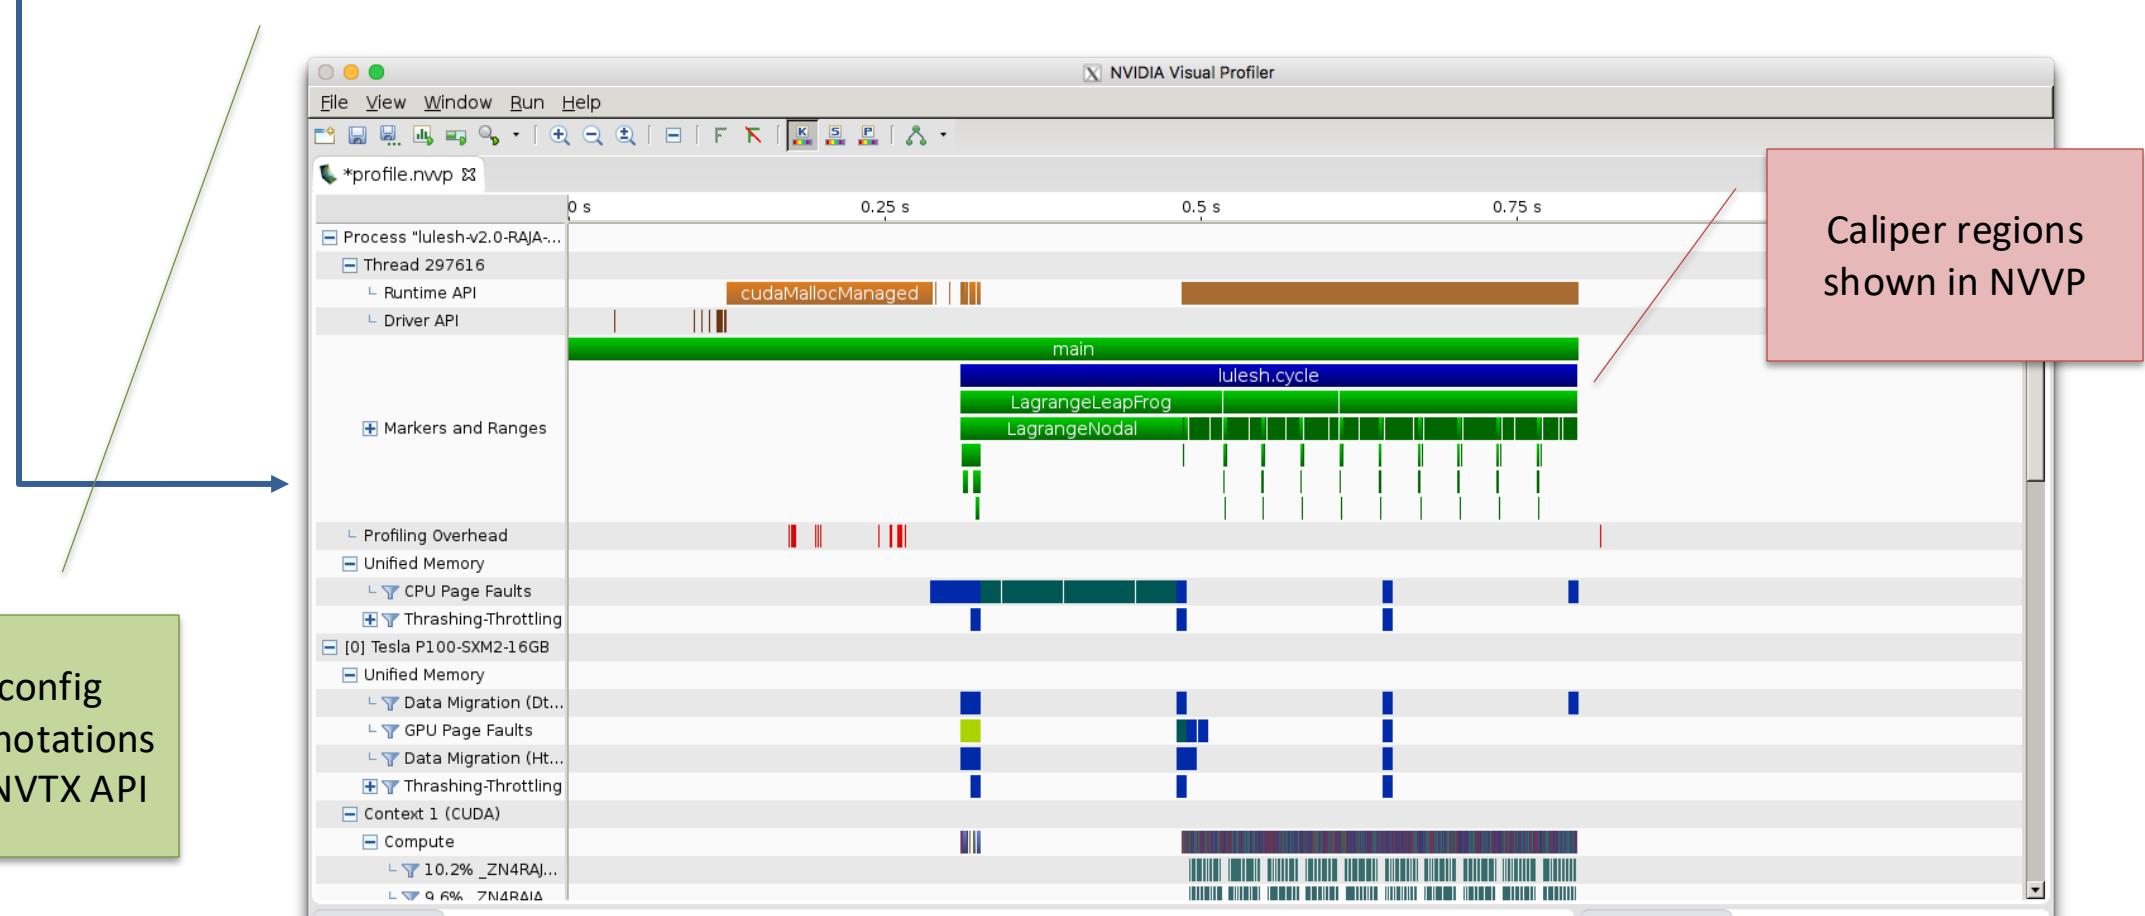
```
$ CALI_CONFIG=event-trace,output=trace.cali ./lulesh2.0  
$ cali2traceevent trace.cali trace.json
```



Caliper event traces can be visualized  
in Chrome trace browser or Perfetto

# Forwarding Annotations to Third-Party Tools

```
$ CALI_CONFIG=nvtx nvprof <nvprof-opt> ./app
```



# Generating Reports with cali-query

```
$ mpirun -n 8 lulesh2.0 -P hatchet-region-profile,profile.mpi
$ cali-query -q "select mpi.rank,sum#sum#time.duration where mpi.function=MPI_Allreduce format tree"
region_profile.cali
```

Path	mpi.rank	time
main		
lulesh.cycle		
TimeIncrement		
MPI_Allreduce		
-	0	0.002174
-	1	0.083150
-	2	0.100689
-	3	0.171060
-	4	0.171347
-	5	0.204563
-	6	0.161819
-	7	0.203488

- Run SQL-like queries on Caliper output with the cali-query tool

# Control Profiling Programmatically: The ConfigManager API

```
#include <caliper/cali.h>
#include <caliper/cali-manager.h>

int main(int argc, char* argv[])
{
    cali::ConfigManager mgr;
    mgr.add(argv[1]);
    if (mgr.error())
        std::cerr << mgr.error_msg() << "\n";

    mgr.start();
    // ...
    mgr.flush();
}
```

- Use ConfigManager to access Caliper's built-in profiling configurations

```
$ ./app runtime-report
```

- Now we can use command-line arguments or other program inputs to enable profiling

# Manual Configuration Allows Custom Analyses

```
cali-query -q "select alloc.label#cuhti.fault.addr as Pool,  
cuhti.uvm.kind as UVM\ Event,  
scale(cuhti.uvm.bytes,1e-6) as MB,  
scale(cuhti.activity.duration,1e-9) as Time  
group by  
prop:nested,alloc.label#cuhti.fault.addr,cuhti.uvm.kind  
where cuhti.uvm.kind format tree" trace.cali
```

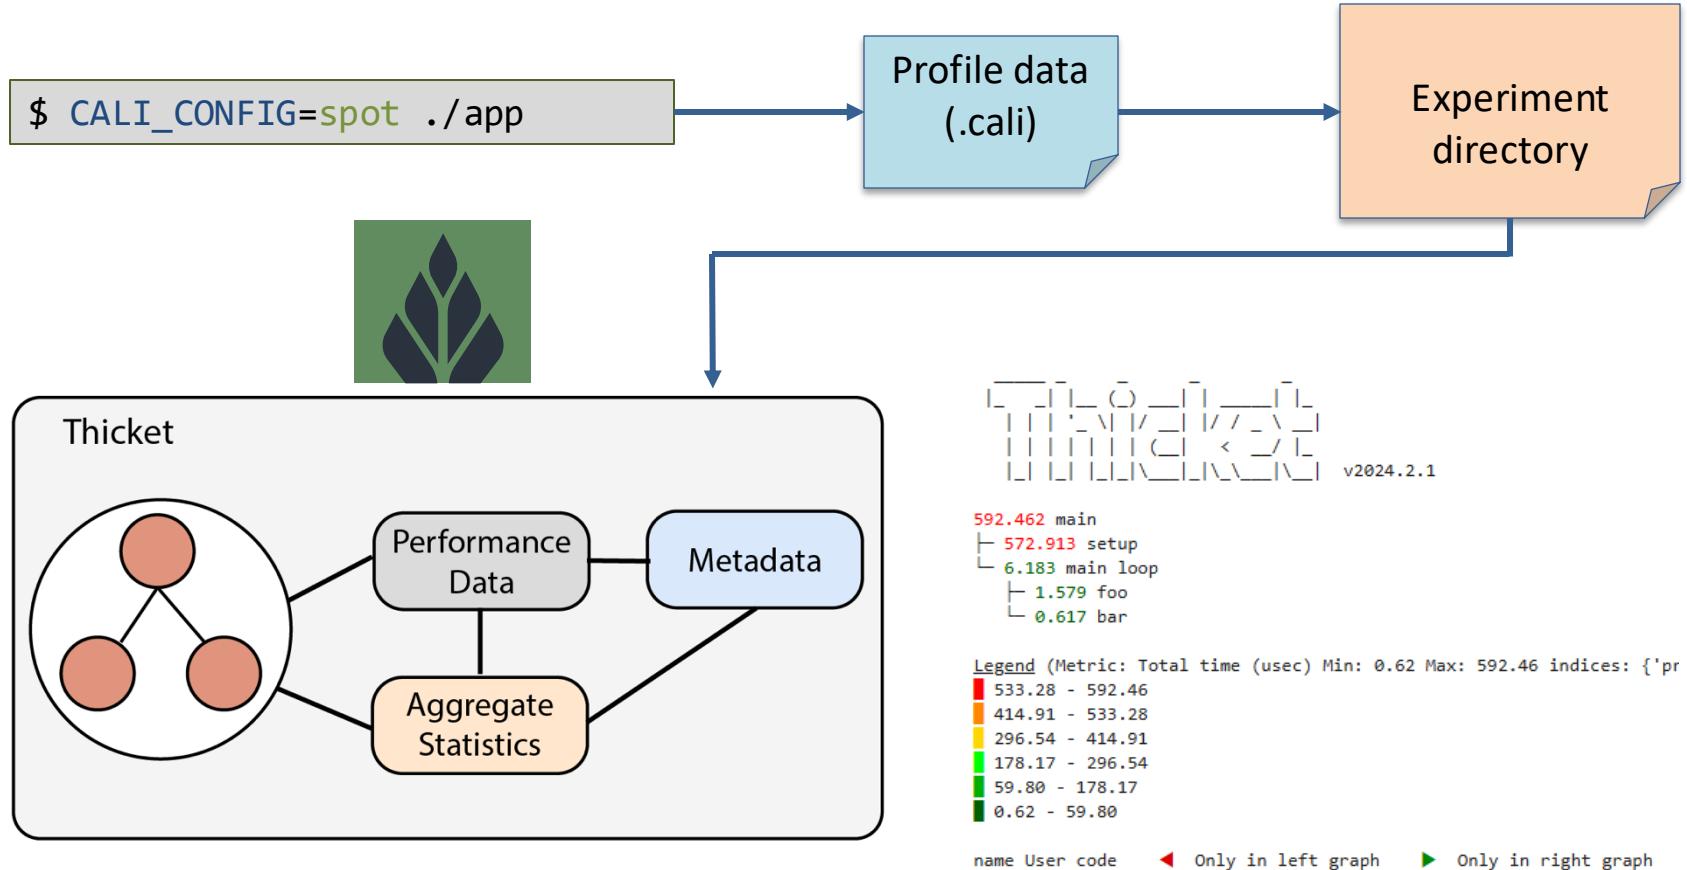
caliper.config

```
CALI_SERVICES_ENABLE=alloc,cuhti,cuhtitrace,mpi,trace,recorder  
CALI_ALLOC_RESOLVE_ADDRESSES=true  
CALI_CUHTI_CALLBACK_DOMAINS=sync  
CALI_CUHTITRACE_ACTIVITIES=uvm  
CALI_CUHTITRACE_CORRELATE_CONTEXT=false  
CALI_CUHTITRACE_FLUSH_ON_SNAPSHOT=true
```

Path				
main				
solve				
TIME_STEPPING				
enforceBC				
CURVI in EnforceBC				
CurviCartIC				
CurviCartIC::PART 3 Pool	UVM Event	MB	Time	
curvilinear4sgwind UM_pool	pagefaults.gpu		2.806946	
curvilinear4sgwind UM_pool	HtoD	7862.747136	0.232238	
curvilinear4sgwind UM_pool_temps	pagefaults.gpu		0.130167	
curvilinear4sgwind UM_pool	DtoH	9986.441216	0.378583	
curvilinear4sgwind UM_pool	pagefaults.cpu			

- Mapping CPU/GPU unified memory transfer events to Umpire memory pools in SW4

# Performance Comparison Studies with Thicket



Thicket: Exploratory Data Analysis in Python for large sets of runs

# Recording Program Metadata with the Adiak Library

TeaLeaf\_CUDA example [C++]

```
#include <adiak.hpp>

adiak::user();
adiak::launchdate();
adiak::jobsizes();

adiak::value("end_step", readInt(input, "end_step"));
adiak::value("halo_depth", readInt(input, "halo_depth"));

if (tl_use_ppcg) {
    adiak::value("solver", "PPCG");
// [...]
```

Use built-in Adiak functions to collect common metadata

Use key:value functions to collect program-specific data

- Use the [Adiak C/C++ library](#) to record program metadata
  - Environment info (user, launchdate, system name, ...)
  - Program configuration (input problem description, problem size, ...)
- Enables performance comparisons across runs. Required for SPOT and Thicket.

# Adiak: Built-in Functions for Common Metadata

```
adiak_user();          /* user name */  
adiak_uid();          /* user id */  
adiak_launchdate();   /* program start time (UNIX timestamp) */  
adiak_executable();   /* executable name */  
adiak_executablepath();/* full executable file path */  
adiak_cmdline();      /* command line parameters */  
adiak_hostname();     /* current host name */  
adiak_clustername();  /* cluster name */  
  
adiak_job_size();     /* MPI job size */  
adiak_hostlist();     /* all host names in this MPI job */  
  
adiak_walltime();     /* wall-clock job runtime */  
adiak_cputime();      /* job cpu runtime */  
adiak_systime();      /* job sys runtime */  
  
adiak_collect_all();  /* everything */
```

- Adiak comes with built-in functions to collect common environment metadata

# Adiak: Recording Custom Key-Value Data in C++

C++

```
#include <adiak.hpp>

vector<int> ints { 1, 2, 3, 4 };
adiak::value("myvec", ints);

adiak::value("myint", 42);
adiak::value("mydouble", 3.14);
adiak::value("mystring", "hi");

adiak::value("mypath", adiak::path("/dev/null"));
adiak::value("compiler", adiak::version("gcc@8.3.0"));
```

- Adiak supports many basic and structured data types
  - Strings, integers, floating point, lists, tuples, sets, ...
- `adiak::value()` records key:value pairs with overloads for many data types

# Adiak: Recording Custom Key-Value Data in C

C

```
#include <adiak.h>

int ints[] = { 1, 2, 3, 4 };
adiak_namevalue("myvec",    adiak_general, NULL, "[%d]", ints, 4);

adiak_namevalue("myint",    adiak_general, NULL, "%d", 42);
adiak_namevalue("mydouble", adiak_general, NULL, "%f", 3.14);
adiak_namevalue("mystring", adiak_general, NULL, "%s", "hi");

adiak_namevalue("mypath",   adiak_general, NULL, "%p", "/dev/null");
adiak_namevalue("compiler", adiak_general, NULL, "%v", "gcc@8.3.0");
```

- In C, `adiak_namevalue()` uses `printf()`-style descriptors to determine data types

---

# Live Demo

## Tutorial Instances: <http://bit.ly/4kGQDlc>

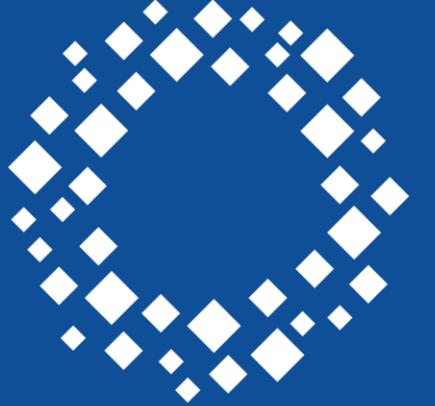
---

- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
  - Pre-installed Thicket, Caliper, and Benchpark and required dependencies
  - Jupyter notebooks
  - Datasets for performance analysis



When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
  - First initial followed by last name (e.g., jdoe)
- PW: hpctutorial2025



# CASC

Center for Applied  
Scientific Computing



**Lawrence Livermore  
National Laboratory**

#### **Disclaimer**

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

---

# Performance Analysis with Thicket I

Stephanie Brink





# Performance Analysis using Thicket

HPDC Tutorial

20 July 2025



Stephanie Brink



LLNL-CFPRES-2202755

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

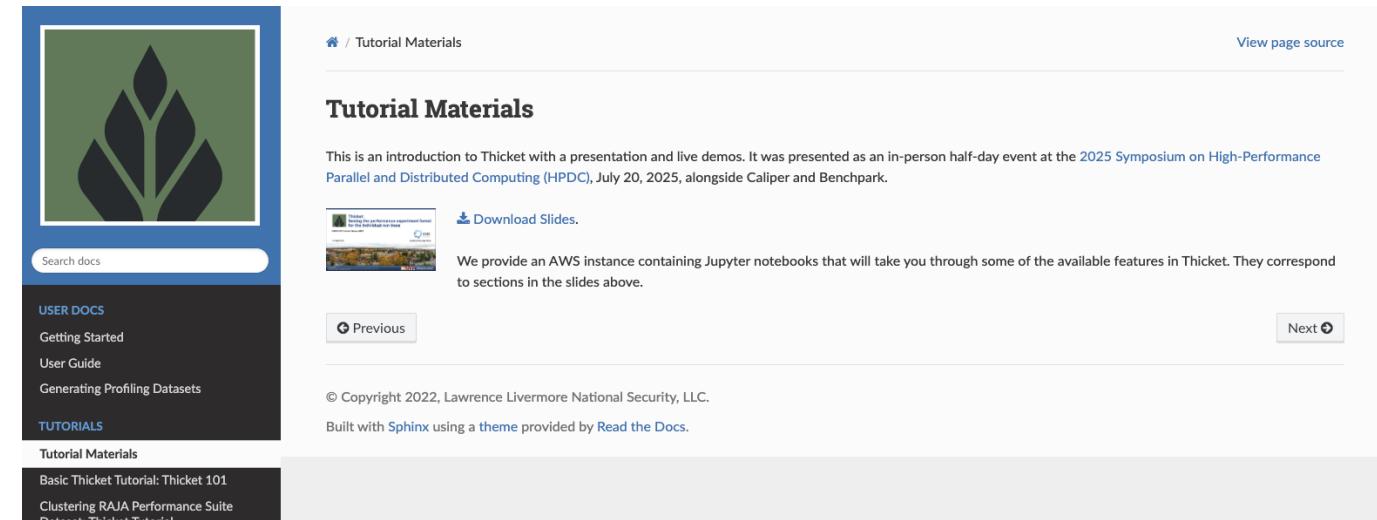


# Tutorial Getting Started

**Presentation slides:** [https://thicket.readthedocs.io/en/latest/tutorial\\_materials.html](https://thicket.readthedocs.io/en/latest/tutorial_materials.html)

## Tutorial Agenda

- Welcome and overview
- Presentation (with slides)
- Hands-on session (AWS instance)



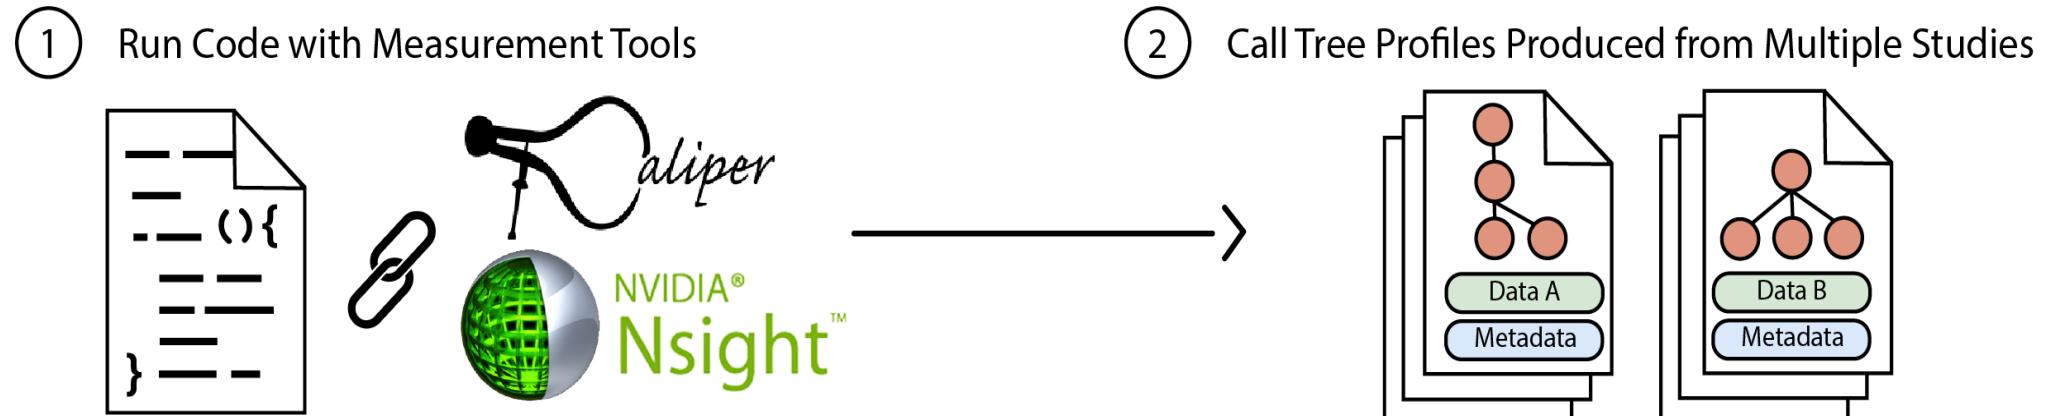
The screenshot shows a documentation page for Thicket. At the top is a navigation bar with a logo of three overlapping blue shapes, a search bar, and links for "USER DOCS" and "TUTORIALS". The main content area has a green header with the text "Tutorial Materials". Below the header, a paragraph describes the page as an introduction to Thicket with a presentation and live demos, mentioning the 2025 Symposium on High-Performance Parallel and Distributed Computing (HPDC). It includes a thumbnail of a presentation slide showing a cluster of servers, a "Download Slides" button, and a note about an AWS instance containing Jupyter notebooks. Navigation buttons for "Previous" and "Next" are at the bottom.

AWS instance included in the tutorial materials provides:

- Pre-installed Thicket and required dependencies
- Jupyter notebooks and associated datasets for analysis

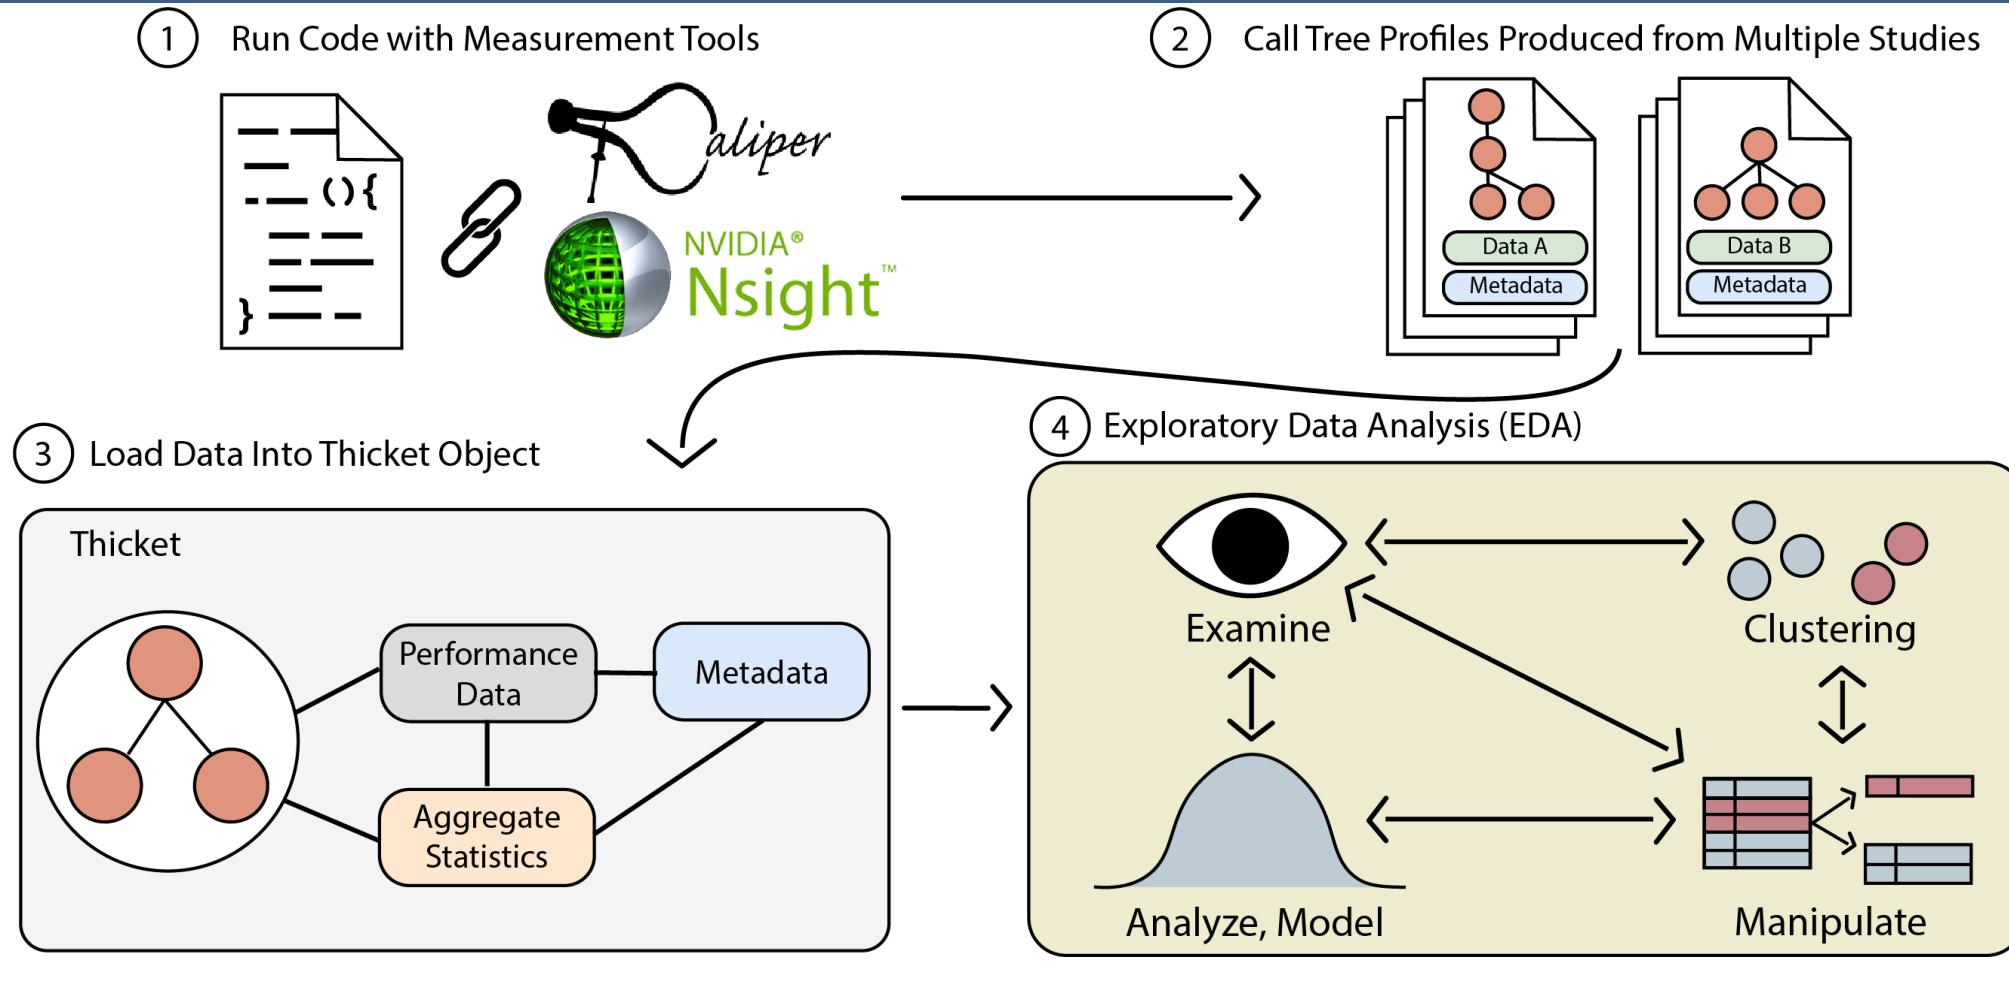
# Challenge: Performance analysis in complex HPC ecosystem

- HPC software and hardware are increasingly complex. Need to understand:
  - Strong scaling and weak scaling of applications
  - Impact of application parameters on performance
  - Impact of choice of compilers and optimization levels
  - Performance on different hardware architectures (e.g., CPUs, GPUs)
  - Different tools to measure different aspects of application performance



Goal: Analyze and visualize performance data from different sources and types

# Our big picture solution for analyzing and visualizing performance data from different sources and type



<https://github.com/LLNL/thicket>

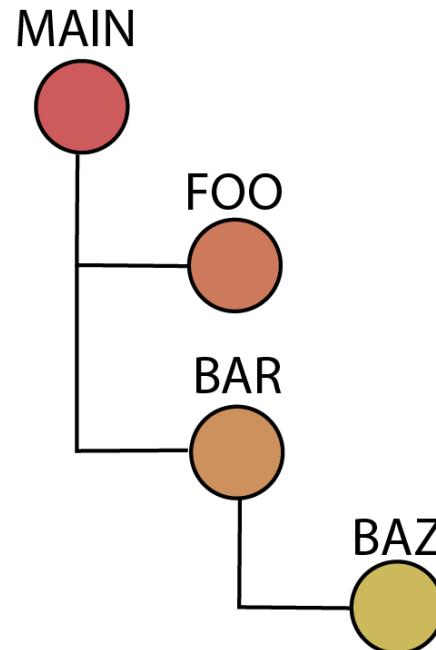
<https://github.com/LLNL/thicket-tutorial>



# What do profiling tools collect per run?



## 1) Call Tree



## 2) Performance data

Node	Cache Misses
MAIN	
FOO	
BAR	
BAZ	

- Time, FLOPS
- Cache misses
- Memory accesses

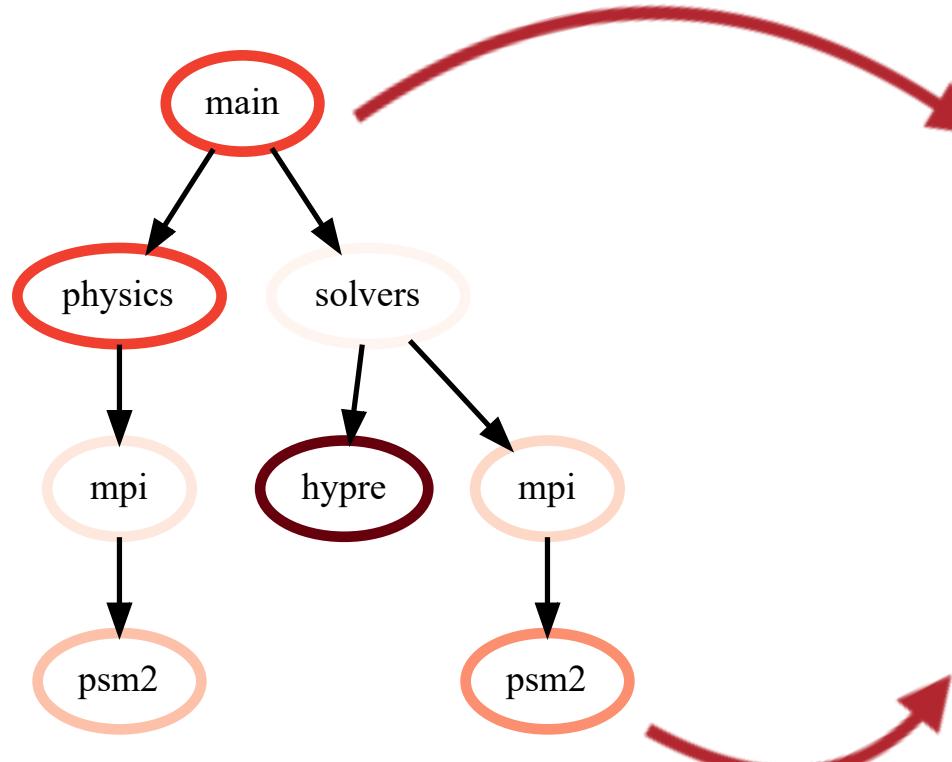
## 3) Metadata per run

User	Platform

- Batch submission (user, launch date)
- Hardware info (platform)
- Build info (compiler versions/flags)
- Runtime info (problem parameters, number of MPI ranks used)



# Thicket builds upon Hatchet's *GraphFrame*: a Graph and a Dataframe



**Graph:** Stores relationships between parents and children

	name	nid	node	time	time (inc)
node					
<b>main</b>	main	0	main	40.0	200.0
<b>physics</b>	physics	1	physics	40.0	60.0
<b>mpi</b>	mpi	2	mpi	5.0	20.0
<b>psm2</b>	psm2	3	psm2	15.0	15.0
<b>solvers</b>	solvers	4	solvers	0.0	100.0
<b>hypre</b>	hypre	5	hypre	65.0	65.0
<b>mpi</b>	mpi	6	mpi	10.0	35.0
<b>psm2</b>	psm2	7	psm2	25.0	25.0

**Pandas Dataframe:** 2D table storing numerical data associated with each node (may be unique per rank, per thread)



<https://github.com/LLNL/hatchet>  
<https://github.com/LLNL/hatchet-tutorial>



# Visualizing Hatchet's GraphFrame components

```
>>> print(gf.tree()) # print graph  
>>> print(gf.dataframe) # print dataframe
```

```
0.000 foo  
|   └ 6.000 bar  
|       |   └ 5.000 baz  
|   └ 0.000 qux  
|       |   └ 5.000 quux  
|           |   └ 10.000 corge  
|           |   └ 15.000 garply  
|               |   └ 1.000 grault  
└ 15.000 waldo  
    |   └ 3.000 fred  
    |       |   └ 5.000 plugh  
    |   └ 15.000 garply
```

Legend (Metric: time)  
■ 13.50 - 15.00  
■ 10.50 - 13.50  
■ 7.50 - 10.50  
■ 4.50 - 7.50  
■ 1.50 - 4.50  
■ 0.00 - 1.50

name User code

◀ Only in left graph

node	name	time	time (inc)
{'name': 'foo'}	foo	0.0	130.0
{'name': 'bar'}	bar	5.0	20.0
{'name': 'baz'}	baz	5.0	5.0
{'name': 'grault'}	grault	10.0	10.0
{'name': 'quux'}	quux	0.0	60.0
{'name': 'quux'}	quux	5.0	60.0
{'name': 'corge'}	corge	10.0	55.0
{'name': 'bar'}	bar	5.0	20.0
{'name': 'baz'}	baz	5.0	5.0
{'name': 'grault'}	grault	10.0	10.0
{'name': 'garply'}	garply	15.0	15.0
{'name': 'grault'}	grault	10.0	10.0



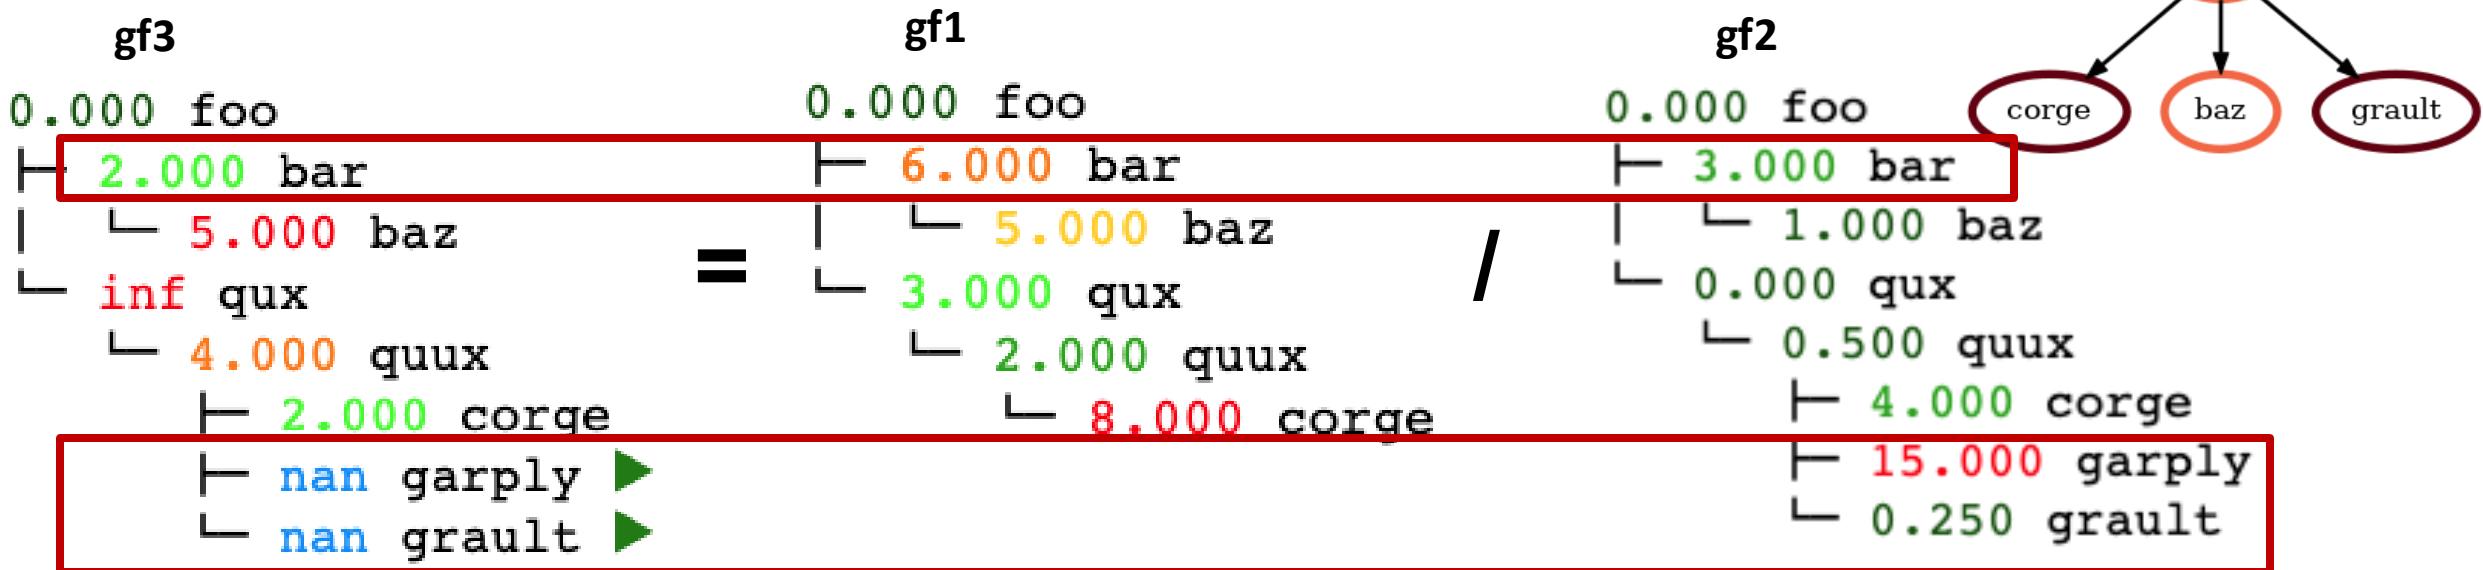
<https://github.com/llnl/benchpark>



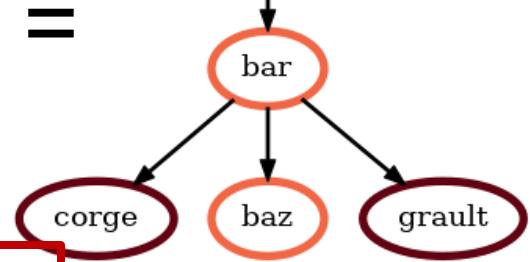
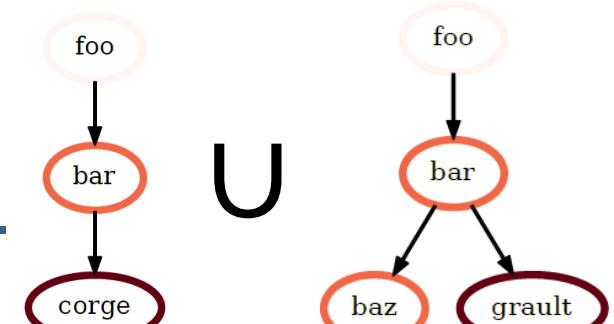
# Compare GraphFrames using division (or add, subtract, multiply)

```
>>> gf3 = gf1 / gf2 # divide graphframes
```

\*First, unify two trees since  
structure is different



```
>>> gf3 = gf1 + gf2 # add graphframes  
>>> gf3 = gf1 - gf2 # subtract graphframes  
>>> gf3 = gf1 * gf2 # multiply graphframes
```



<https://github.com/llnl/benchpark>



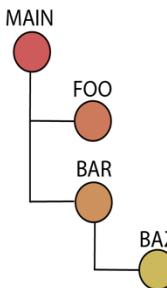


# Use Thicket to *compose* performance profiles in Python

P1

Metadata

User	Platform



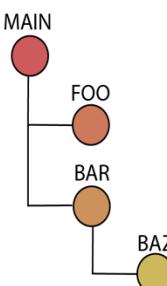
Performance metrics

Node	Cache Misses
MAIN	24
FOO	
BAR	
BAZ	

P2

Metadata

User	Platform

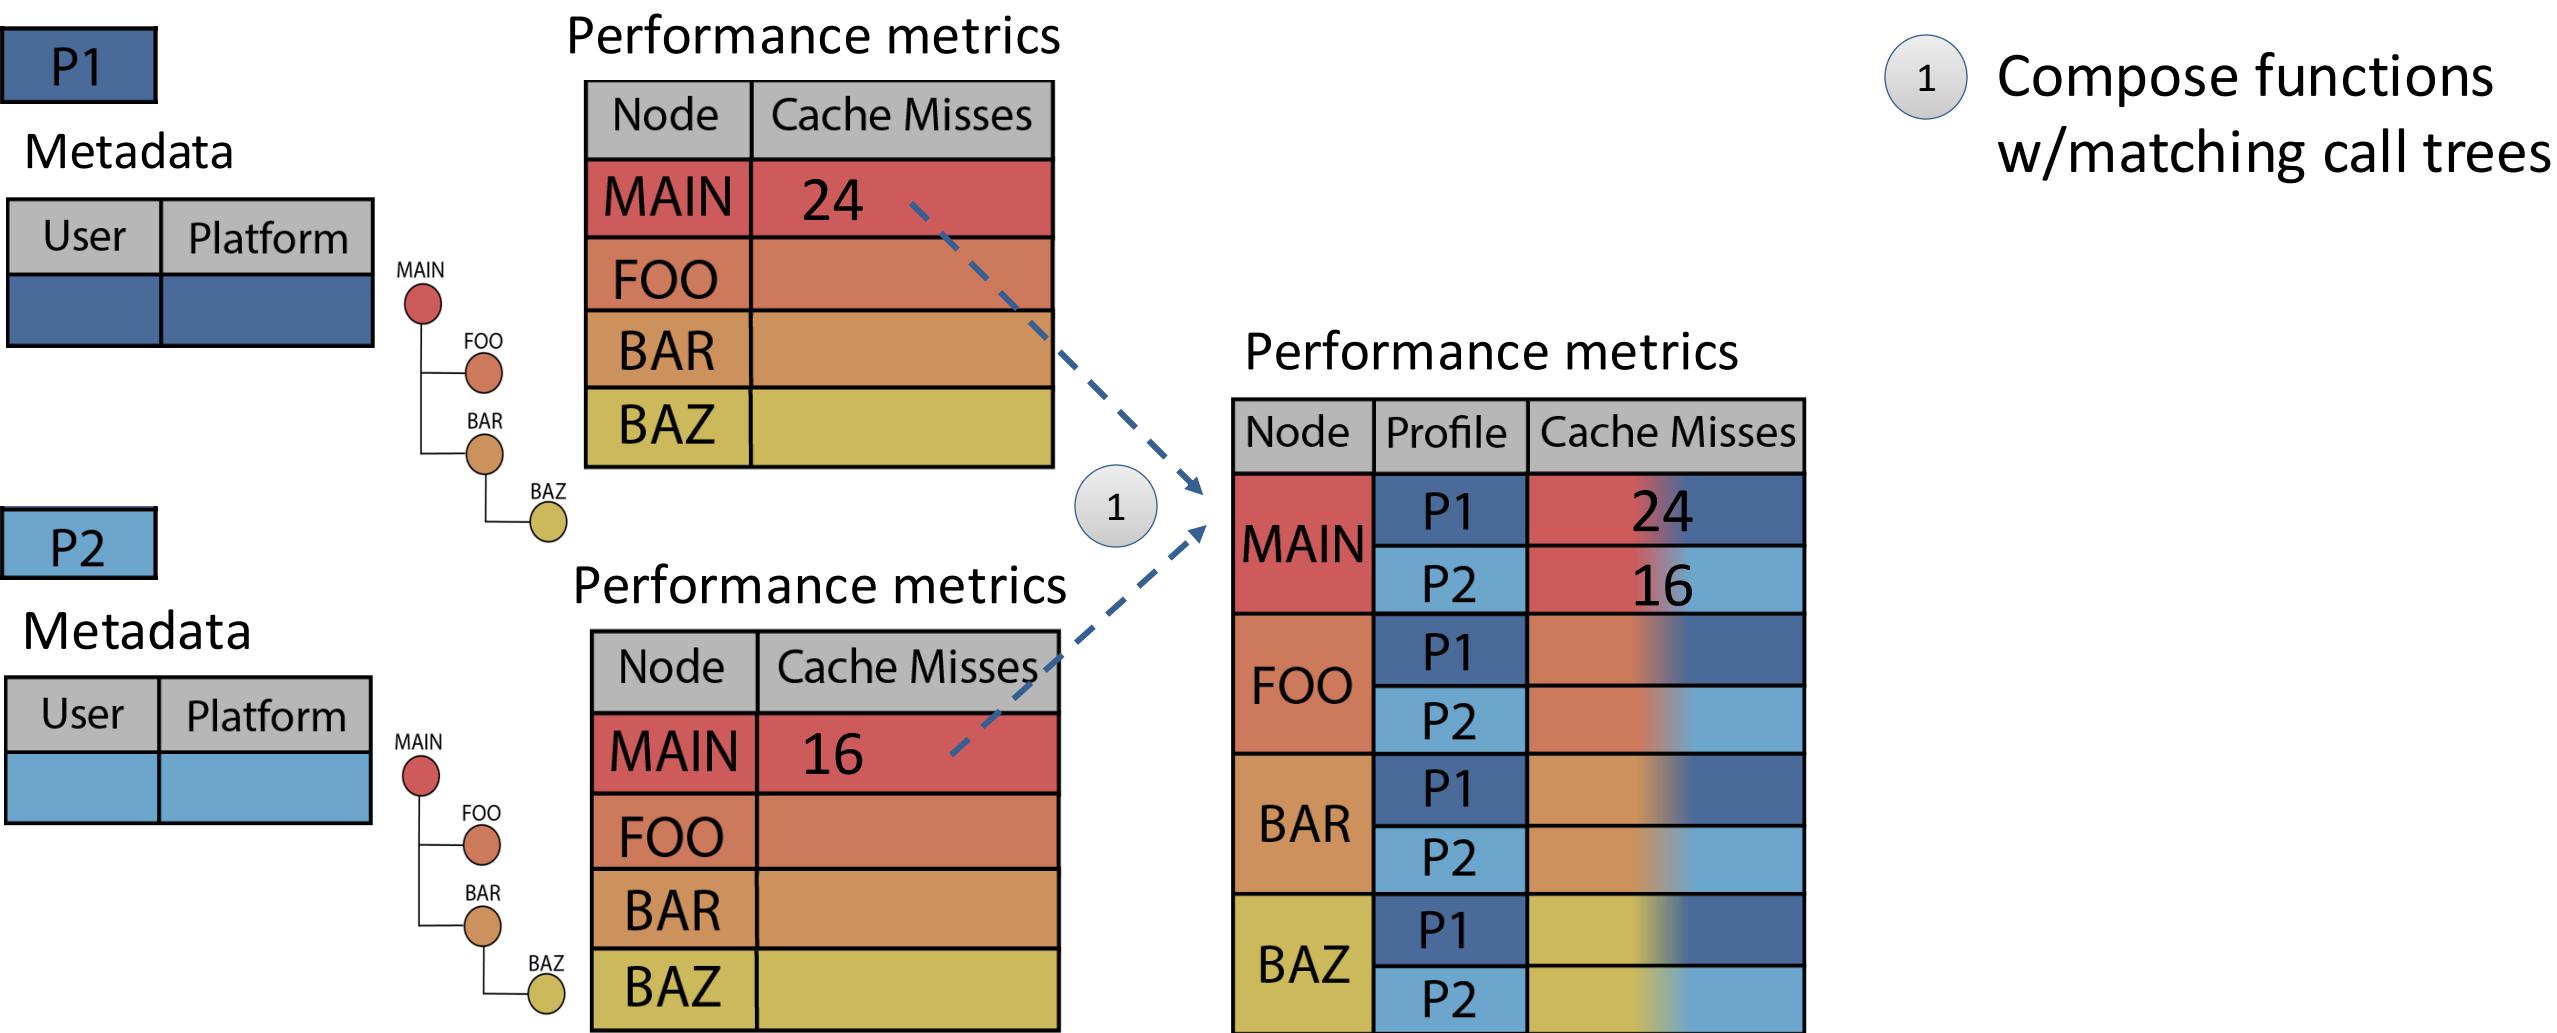


Performance metrics

Node	Cache Misses
MAIN	16
FOO	
BAR	
BAZ	

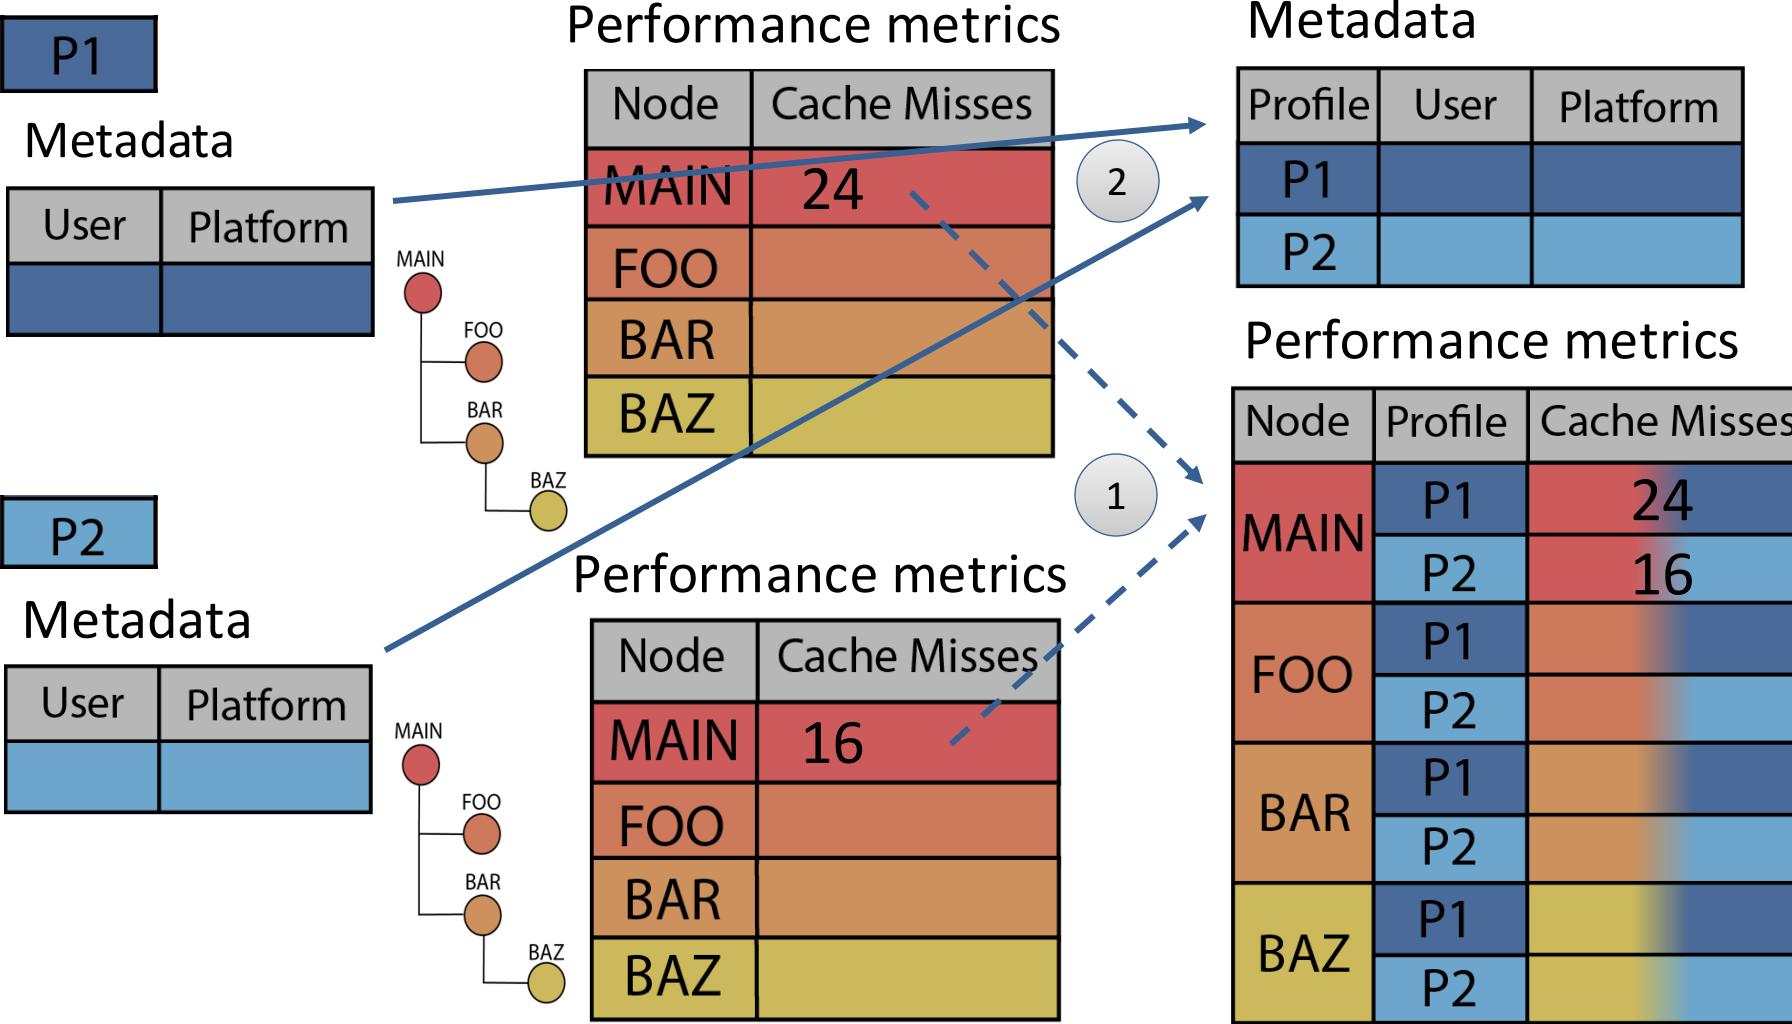


# Use Thicket to *compose* performance profiles in Python





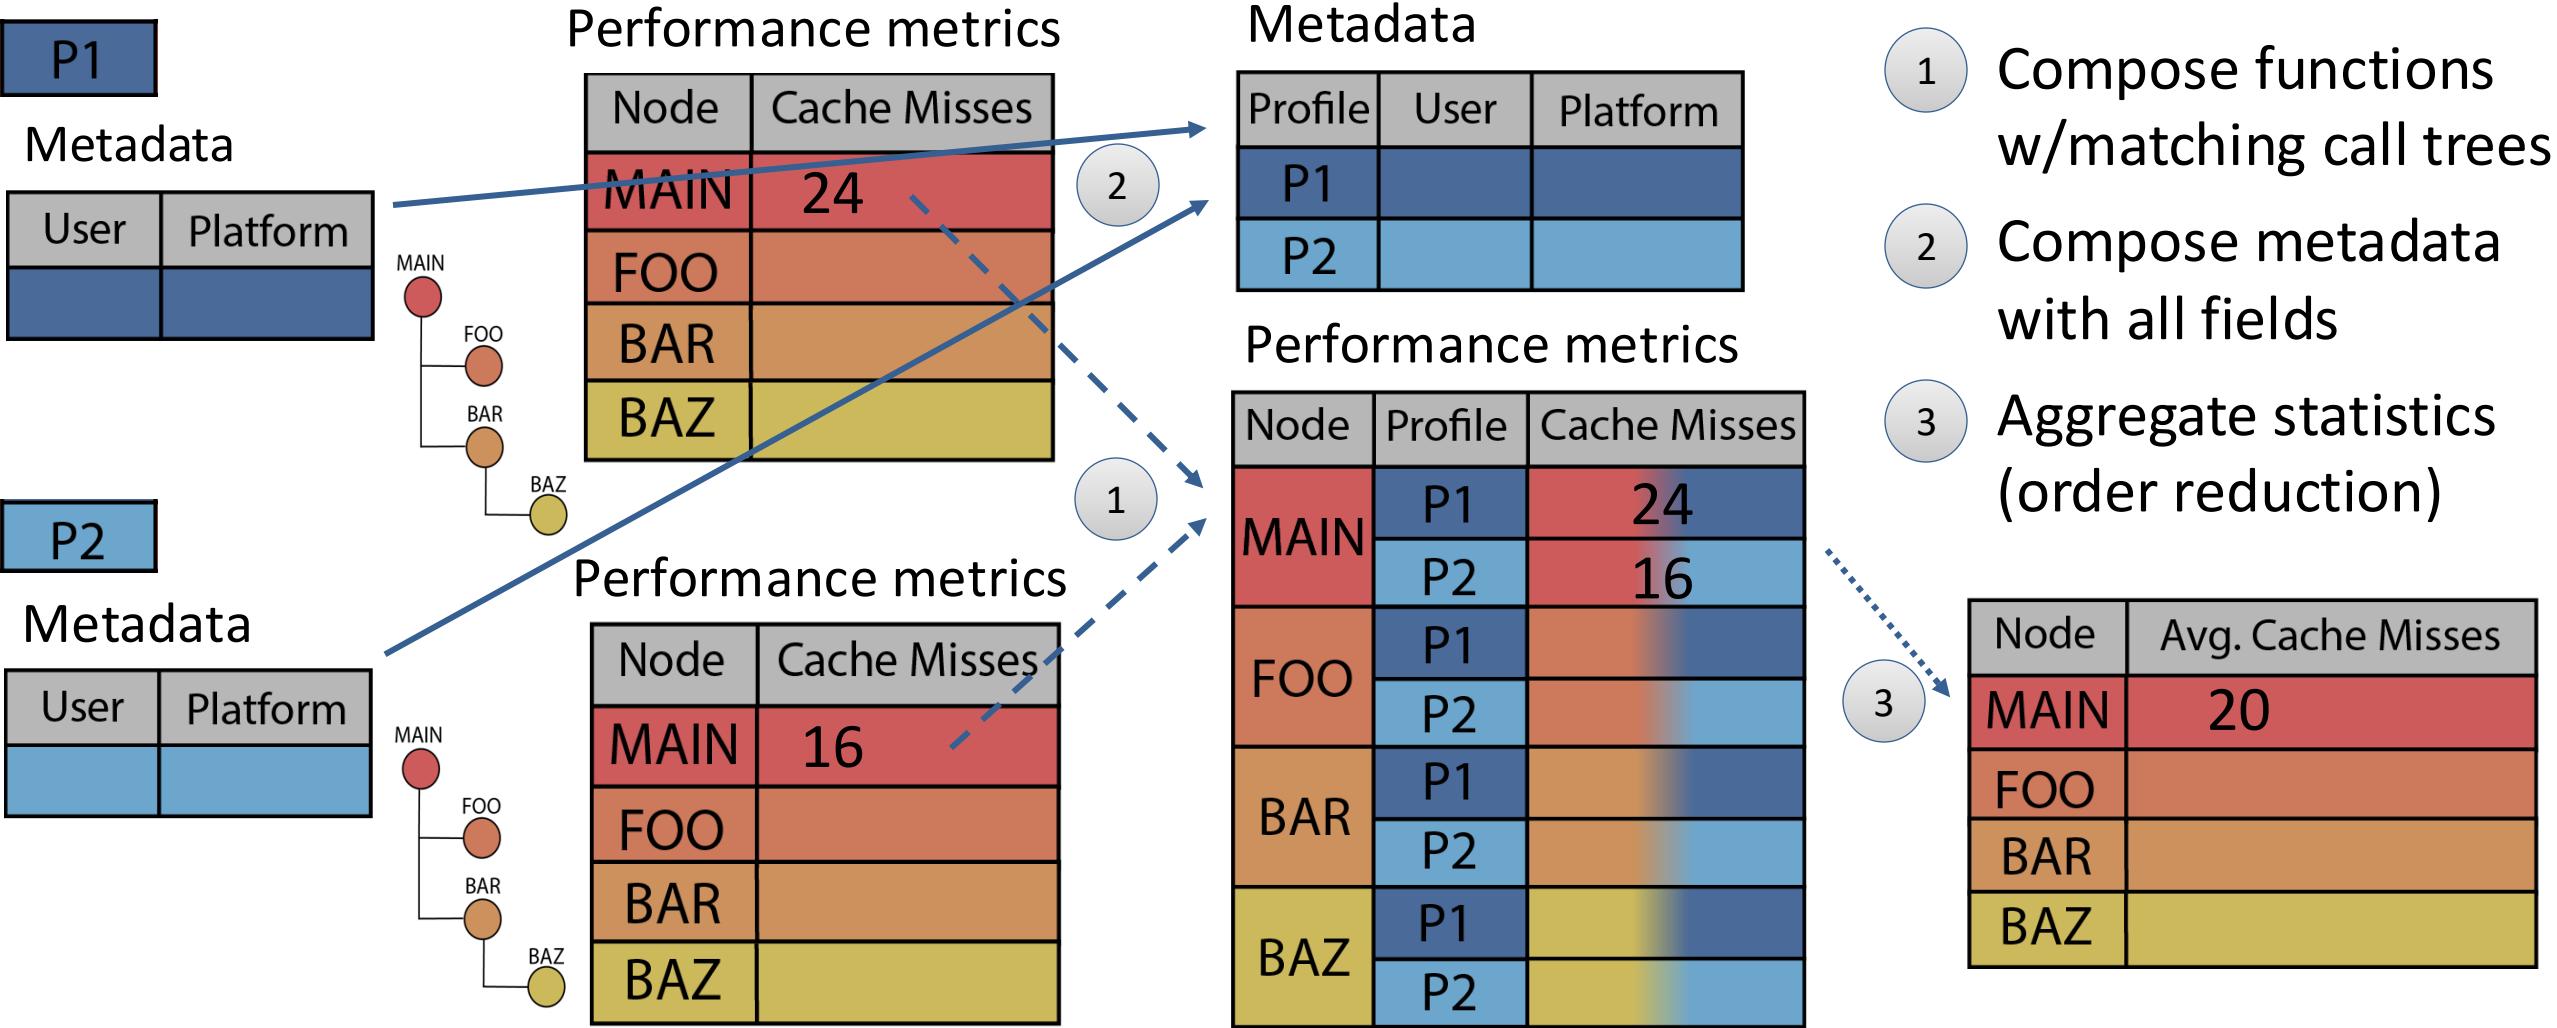
# Use Thicket to *compose* performance profiles in Python



- 1 Compose functions w/matching call trees
- 2 Compose metadata with all fields



# Use Thicket to *compose* performance profiles in Python



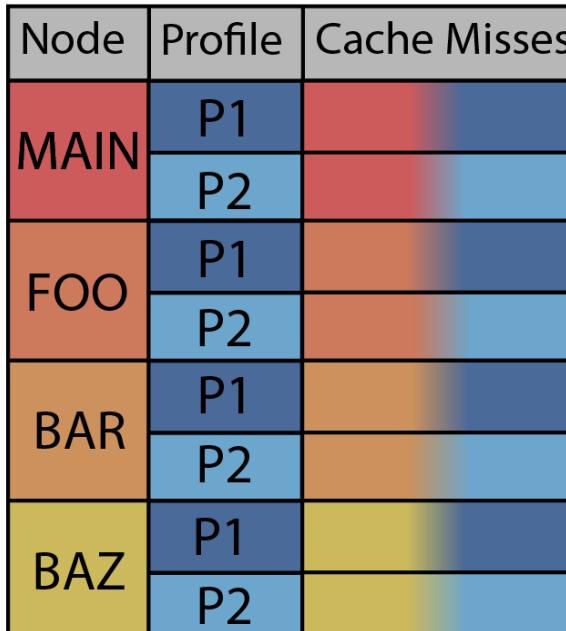


# Thicket components are *interconnected*

Metadata

Profile	User	Platform
P1	Jon	lassen
P2	Bob	lassen

Performance metrics



Filter on metadata:  
platform=="lassen" &&  
user=="Bob"

Filtered Metadata

Profile	User	Platform
P2	Bob	lassen

Filtered Performance metrics

Node	Profile	Cache Misses
MAIN	P2	High (Red)
FOO	P2	Medium (Orange)
BAR	P2	Low (Blue)
BAZ	P2	Very Low (Yellow)

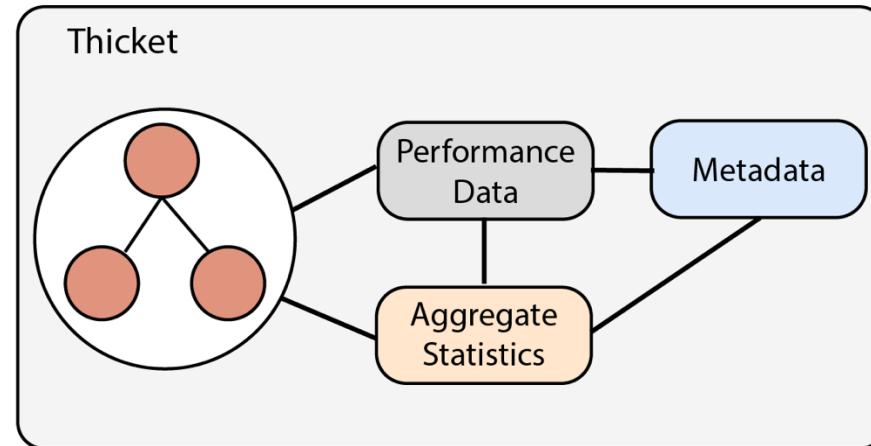
Metadata fields useful for understanding  
and manipulating thicket object!





# Thicket enables exploratory data analysis of multi-run data

## ③ Load Data Into Thicket Object

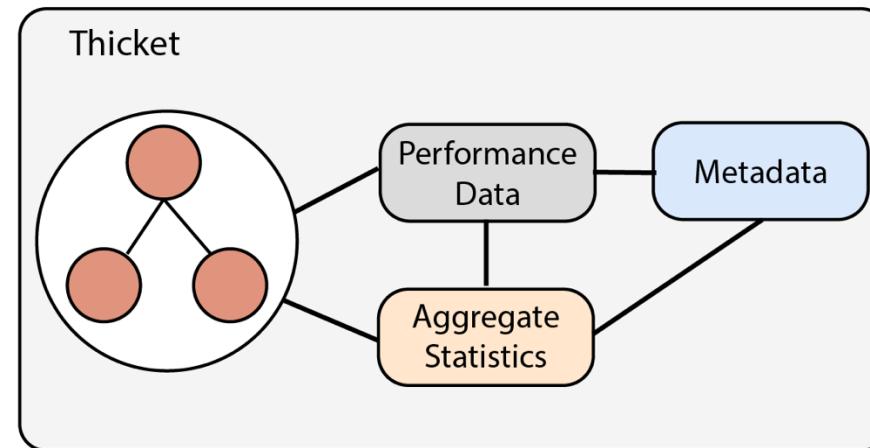


- Compose data from diff. sources and types
  - Different scaling (e.g., strong, weak)
  - Different application parameters
  - Different compilers and optimization levels
  - Different hardware types (e.g., CPUs, GPUs)
  - Different performance tools

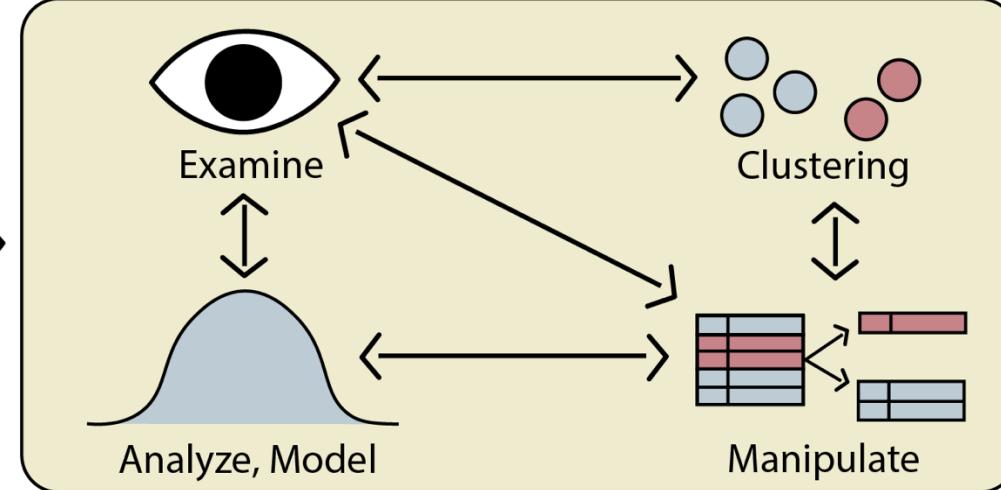


# Thicket enables exploratory data analysis of multi-run data

## 3 Load Data Into Thicket Object



## 4 Exploratory Data Analysis (EDA)



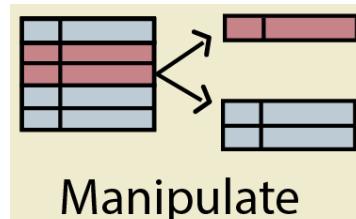
- Compose data from diff. sources and types
  - Different scaling (e.g., strong, weak)
  - Different application parameters
  - Different compilers and optimization levels
  - Different hardware types (e.g., CPUs, GPUs)
  - Different performance tools
- Perform analysis on the thicket of runs
  - Manipulate the set of data
  - Visualize the dataset
  - Perform analysis on the data
  - Model data
  - Leverage third-party tools in the Python ecosystem



- Open-source suite of loop-based kernels commonly found in HPC applications showcasing performance of different programming models on different hardware
- 560 runs/profiles:
  - 2 clusters (CPU, CPU+GPU)
  - 4 problem sizes
  - 3 compilers, 4 optimizations
  - 3 programming models (sequential, OpenMP, CUDA)
  - 3 performance tools (Caliper, PAPI, Nsight Compute)

<http://github.com/llnl/rajaperf>

cluster	systype build	problem size	compiler	compiler optimizations	omp num threads	cuda compiler	block sizes	RAJA variant	#profiles
0	quartz	toss_3_x86_64_ib	[1M, 2M, 4M, 8M]	clang++-9.0.0	[-O0, -O1, -O2, -O3]	1	N/A	N/A	Sequential 160
1	quartz	toss_3_x86_64_ib	[1M, 2M, 4M, 8M]	g++-8.3.1	[-O0, -O1, -O2, -O3]	1	N/A	N/A	Sequential 160
2	quartz	toss_3_x86_64_ib	[1M, 2M, 4M, 8M]	clang++-9.0.0	-O0	72	N/A	N/A	OpenMP 40
3	quartz	toss_3_x86_64_ib	[1M, 2M, 4M, 8M]	g++-8.3.1	-O0	72	N/A	N/A	OpenMP 40
4	lassen	blueos_3_ppc64le_ib_p9	[1M, 2M, 4M, 8M]	xlc++_r-16.1.1.12	-O0	1 nvcc-11.2.152	[128, 256, 512, 1024]	CUDA	160



# Use Thicket to *compose* multi-platform, multi-tool data

Thicket object composed of 2 profiles run on CPU

		node	problem_size	time (exc)	Reps	Retiring	Backend bound
Apps_NODAL_ACCUMULATION_3D	1M	0.204583	100	0.144928	0.783786		
		0.795511	100	0.139002	0.788017		
	1M	0.067061	100	0.402238	0.510525		
	4M	0.241508	100	0.400775	0.515976		

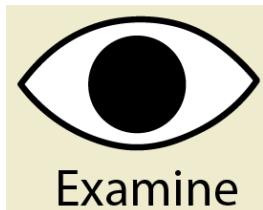
Thicket object composed of 2 profiles run on GPU

		node	problem_size	time (gpu)	gpu_compute_memory_throughput	gpu_dram_throughput	sm_throughput
Apps_NODAL_ACCUMULATION_3D	1M	0.007478		70.689752	46.724767	7.330745	
		0.026951		74.275834	51.257993	7.688628	
	1M	0.006028		81.012826	67.751194	35.676942	
	4M	0.021422		91.929933	70.122011	35.386470	

CPU

GPU

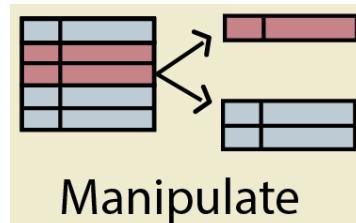
		node	problem_size	time (exc)	Reps	Retiring	Backend bound	time (gpu)	gpu_compute_memory_throughput	gpu_dram_throughput	sm__throughput
Apps_NODAL_ACCUMULATION_3D	1M	0.204583	100	0.144928	0.783786	0.007478		70.689752	46.724767	7.330745	
		0.795511	100	0.139002	0.788017	0.026951		74.275834	51.257993	7.688628	
	1M	0.067061	100	0.402238	0.510525	0.006028		81.012826	67.751194	35.676942	
	4M	0.241508	100	0.400775	0.515976	0.021422		91.929933	70.122011	35.386470	



- Dataset: 4 types of profiles side-by-side to compare CPU to GPU performance
  - 1 Basic CPU metrics from Caliper
  - 2 Top-down metrics from Caliper/PAPI
  - 3 GPU runtime from Caliper
  - 4 GPU metrics from Nsight Compute
- Examples of analysis:
  - Compute CPU/GPU speedup
  - Correlate memory and compute usage on the CPU vs. GPU

Node	Problem size	CPU			CPU top-down		GPU		GPU Nsight Compute				speedup
		time (exc)	Bytes/Rep	Flops/Rep	Retiring	Backend bound	time (gpu)	gpu_compute_memory_throughput	gpu_dram_throughput	sm_throughput	sm_warps_active		
Apps_VOL3D	8M	0.498815	282109496	632421288	0.377843	0.540604	0.040761		93.742058	72.140428	36.206767	54.459589	12.237556
Lcals_HYDRO_1D	8M	2.077556	201326600	41943040	0.032965	0.909545	0.242928		92.944968	92.944968	6.595714	95.266148	8.552147

Derived



# Manipulate: Filter using call path query

```
0.001 Base_CUDA
└ 0.000 Algorithm
  └ 0.000 Algorithm_MEMORY
    └ 0.002 Algorithm_MEMORY.block_128
      └ 0.009 Algorithm_MEMORY.block_256
        └ 0.006 Algorithm_MEMORY.library
  └ 0.000 Algorithm_MEMSET
    └ 0.001 Algorithm_MEMSET.block_128
      └ 0.004 Algorithm_MEMSET.block_256
        └ 0.003 Algorithm_MEMSET.library
  └ 0.000 Algorithm_REDUCE_SUM
    └ 0.003 Algorithm_REDUCE_SUM.block_128
      └ 0.004 Algorithm_REDUCE_SUM.block_256
        └ 0.002 Algorithm_REDUCE_SUM.cub
  └ 0.000 Algorithm_SCAN
    └ 0.006 Algorithm_SCAN.default
```

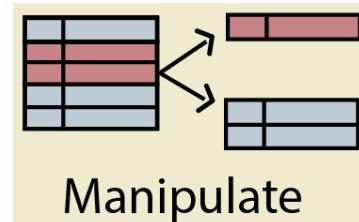
Input call tree

Filter on call path:  
(1) Node named  
“Base\_CUDA”

```
0.001 Base_CUDA
```

Output call tree

I Lumsden et al. “Enabling Call Path Querying in Hatchet to Identify Performance Bottlenecks in Scientific Applications”, e-Science 2022



# Manipulate: Filter using call path query

```
0.001 Base_CUDA
└ 0.000 Algorithm
  └ 0.000 Algorithm_MEMCPY
    └ 0.002 Algorithm_MEMCPY.block_128
    └ 0.009 Algorithm_MEMCPY.block_256
    └ 0.006 Algorithm_MEMCPY.library
  └ 0.000 Algorithm_MEMSET
    └ 0.001 Algorithm_MEMSET.block_128
    └ 0.004 Algorithm_MEMSET.block_256
    └ 0.003 Algorithm_MEMSET.library
  └ 0.000 Algorithm_REDUCE_SUM
    └ 0.003 Algorithm_REDUCE_SUM.block_128
    └ 0.004 Algorithm_REDUCE_SUM.block_256
    └ 0.002 Algorithm_REDUCE_SUM.cub
  └ 0.000 Algorithm_SCAN
    └ 0.006 Algorithm_SCAN.default
```

Input call tree

- Filter on call path:
- (1) Node named “Base\_CUDA”
  - (2) Node with “block\_128” in name (and any nodes in between)

```
0.001 Base_CUDA
└ 0.000 Algorithm
  └ 0.000 Algorithm_MEMCPY
    └ 0.002 Algorithm_MEMCPY.block_128
  └ 0.000 Algorithm_MEMSET
    └ 0.001 Algorithm_MEMSET.block_128
  └ 0.000 Algorithm_REDUCE_SUM
    └ 0.003 Algorithm_REDUCE_SUM.block_128
```

Output call tree

I Lumsden et al. “Enabling Call Path Querying in Hatchet to Identify Performance Bottlenecks in Scientific Applications”, e-Science 2022

---

# 30 Min Break

See you at 3:30pm

Tutorial: Reproducible Performance Measurement and Analysis for High-  
Performance Computing Applications

---

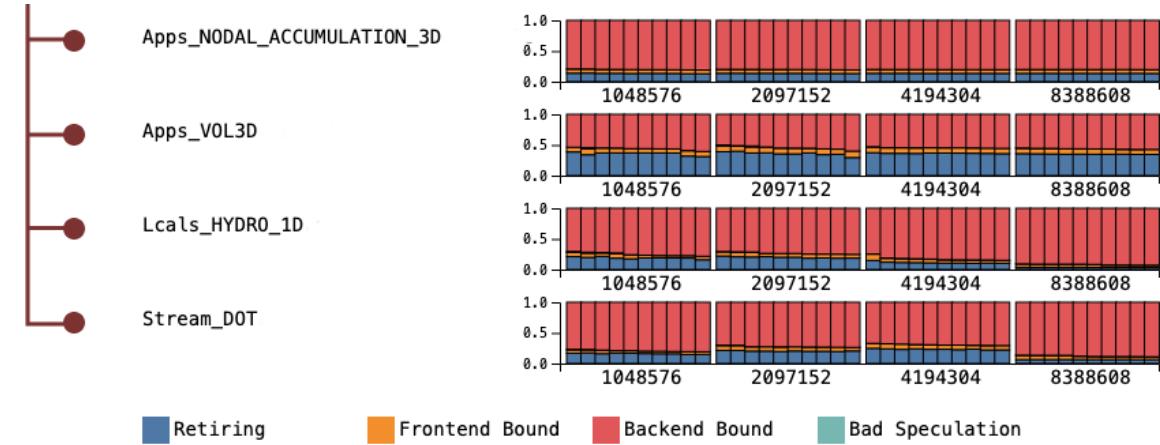
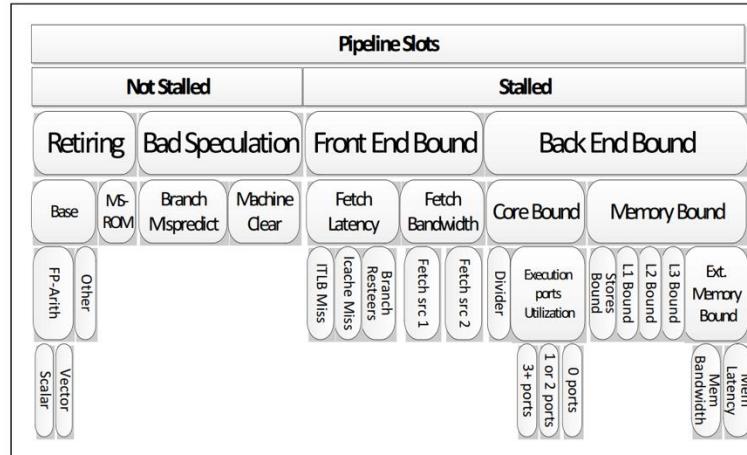
# Performance Analysis with Thicket II

Stephanie Brink





# Visualize: Intel CPU top-down analysis

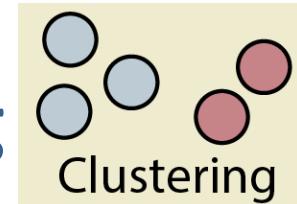


- *Top-down analysis* uses HW counters in a hierarchy to identify bottlenecks\*
- Use Caliper's top-down module to derive top-down metrics for call-tree regions

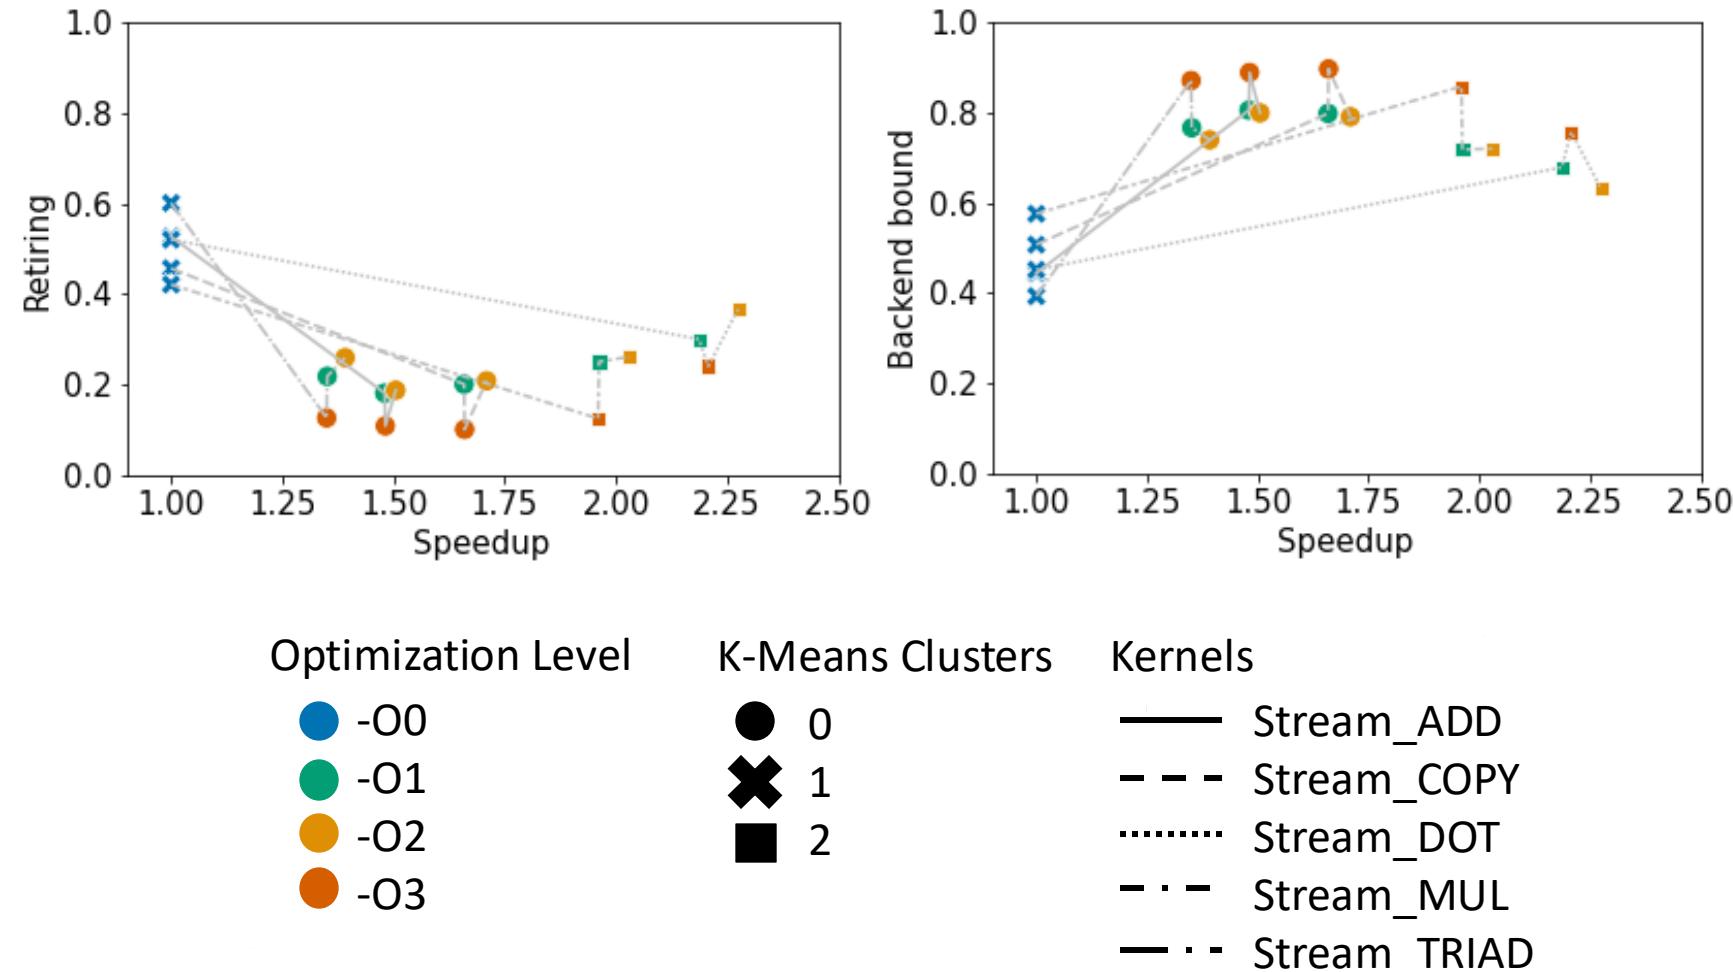
- Thicket's *tree+table* visualization shows top-down metrics as stacked bar charts, each bar is a profile
  - Apps\_VOL3D has the highest retiring rates
  - Lcals\_HYDRO and Stream\_DOT become more backend bound as problem size grows

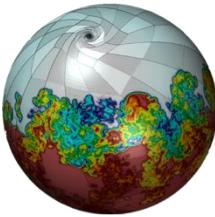
\* Yasin, A.: A Top-Down Method for Performance Analysis and Counters Architecture. In: 2014 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS). pp. 35–44. IEEE, CA, USA (Mar 2014).

# Use third-party Python libraries, e.g., Scikit-learn clustering



1. Select data of interest
  - Filter 8M problem size
  - Use query language to extract all implementations of the Stream kernel
2. (optional) Normalize data
3. Apply scikit-learn clustering to top-down analysis metrics of runs with different compiler optimization levels

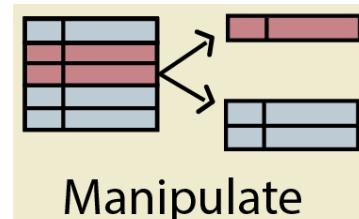




## Case Study 2: MARBL multi-physics code

- MARBL is a next-generation multi-physics code developed at LLNL
- 60 runs/profiles:
  - 2 clusters (rztopaz, AWS ParallelCluster)
  - 2 MPI libraries (impi, openmpi)
  - 6 node/rank counts
  - 5 repeat runs per config

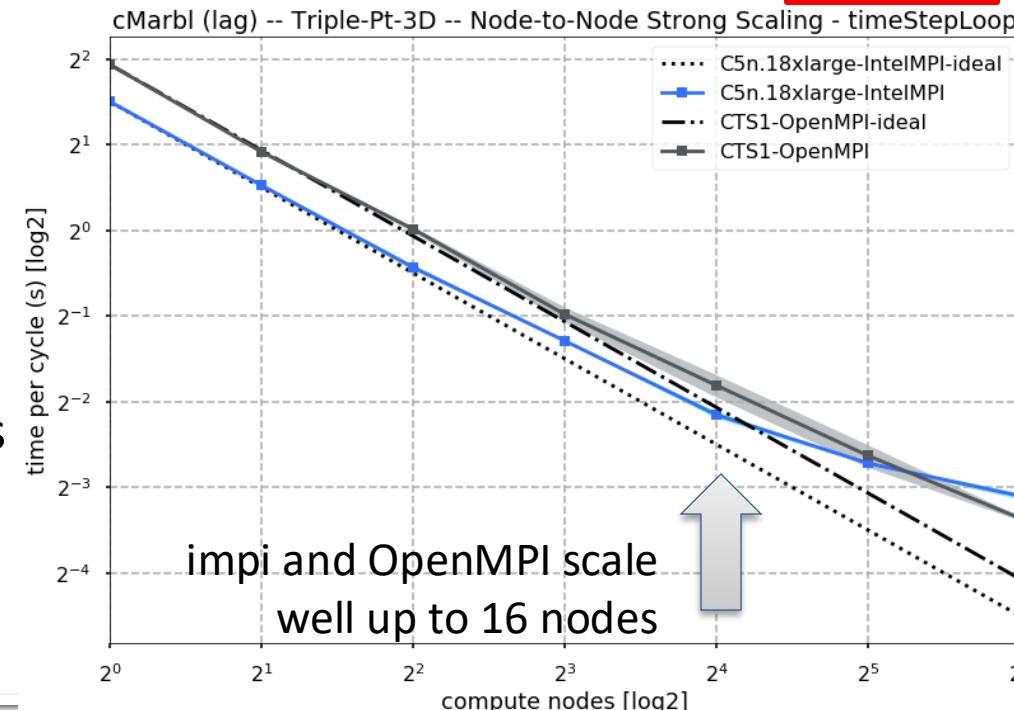
cluster	ccompiler	mpi	version	numhosts	mpi.world.size	#profiles
0 ip----	/usr/tce/packages/clang/clang-9.0.0	impi	v1.1.0-203-gcb0efb3	[1, 2, 4, 8, 16, 32]	[36, 72, 144, 288, 576, 1152]	30
1 rztopaz	/usr/tce/packages/clang/clang-9.0.0	openmpi	v1.1.0-201-g891eaf1	[1, 2, 4, 8, 16, 32]	[36, 72, 144, 288, 576, 1152]	30



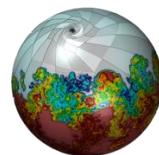
# Manipulate: Compute noise and scaling

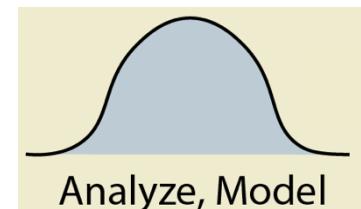
node	profile	Total time	name	mpi.world.size
{'name': 'main', 'type': 'function'}	-8554409769265002864	58036.664552	main	144
	-7335101512240609798	55318.808836	main	36
	-6029692086108825020	156984.246813	main	2304
	-5606382734792961361	64122.371533	main	288
	-4058809097109060732	155040.998627	main	2304
	-3193575964635936033	71010.504038	main	576
	-2978339073585311581	55910.708449	main	72
	-2939704488254773514	157934.204076	main	2304
	-2771797711381234985	56893.512948	main	144
	-2638513839856695106	97432.260966	main	1152

node	profile	Total time	name	mpi.world.size
	{'name': 'main', 'type': 'function'}		-7335101512240609798	55318.808836
			-843517585394879415	55110.656885
			7720382918482619866	55155.581578
			8293335926964337960	55139.134916
			8335957980556391465	55013.682102



1. Use groupby(MPI.world.size) to generate unique subsets of data which are repeated runs; compute noise
2. Compose runs on different platforms and at different scales
3. Generate strong scaling plot with matplotlib
  - Deviation shown in shaded region, dots are average of 5 runs



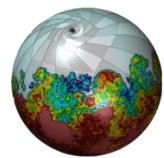
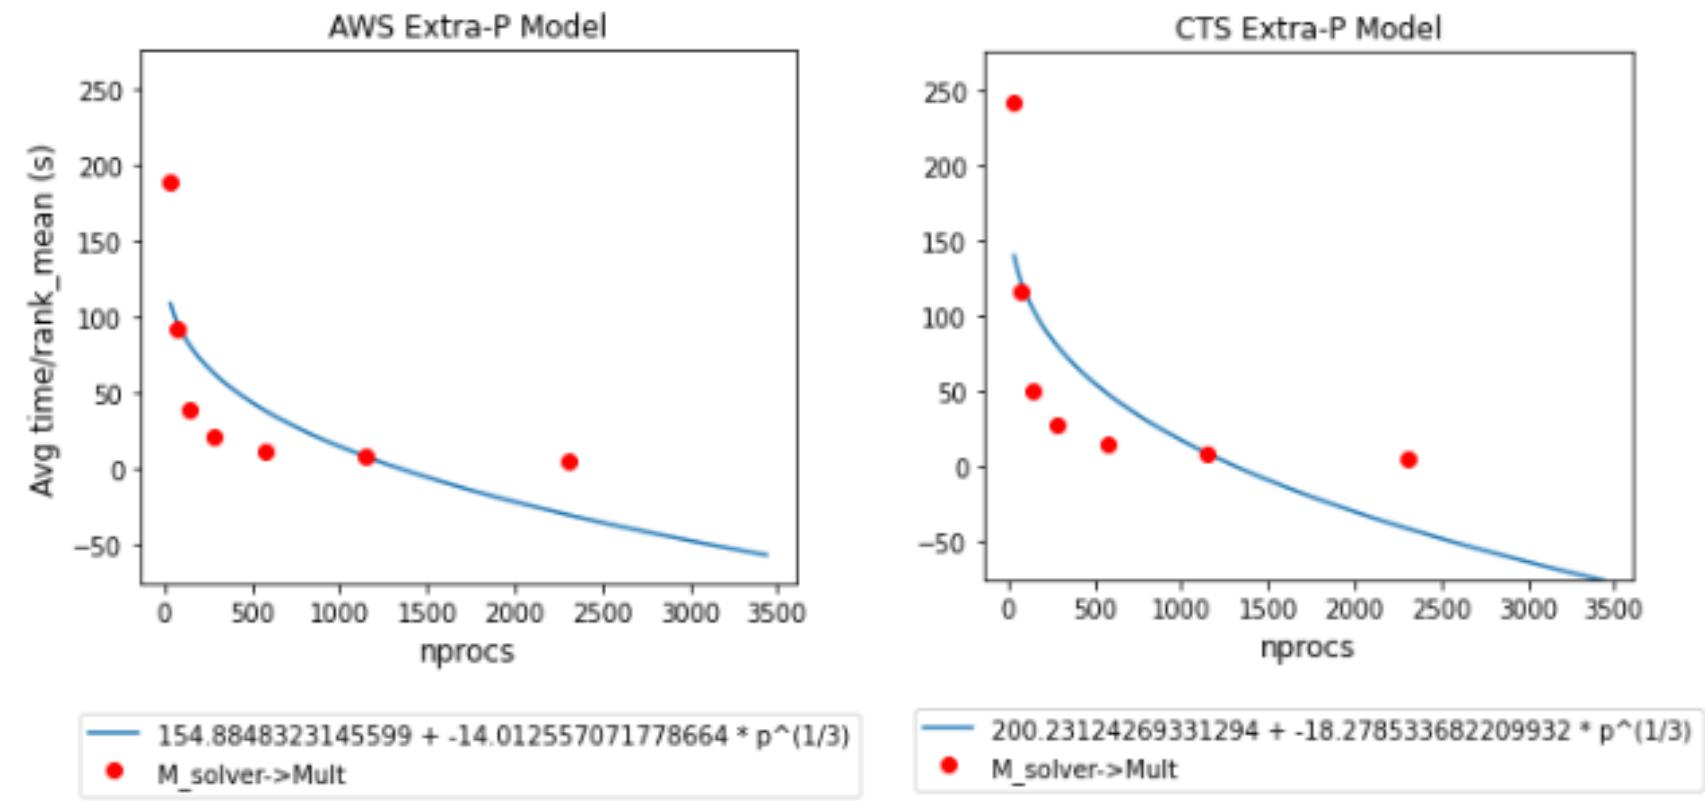


# Model: Use third-party Python library, Extra-P

Extra-P derives an analytical performance model from an ensemble of profiles covering one or more modeling parameters

<http://github.com/extra-p/extrap>

- Select functions of interest
- Call Extra-P to model scaling on different hardware types





# Visualize metadata with parallel coordinates plot

- Thicket's interactive parallel coordinates plot shows relationships between metadata variables, and between metadata and performance data

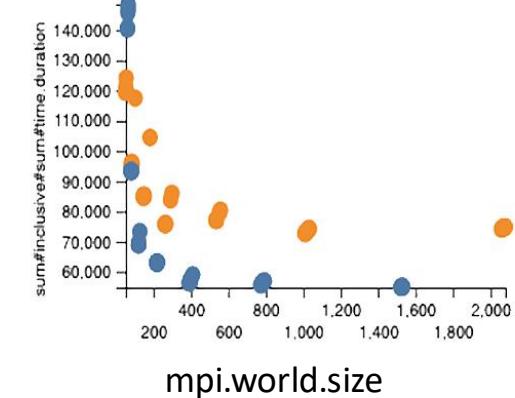
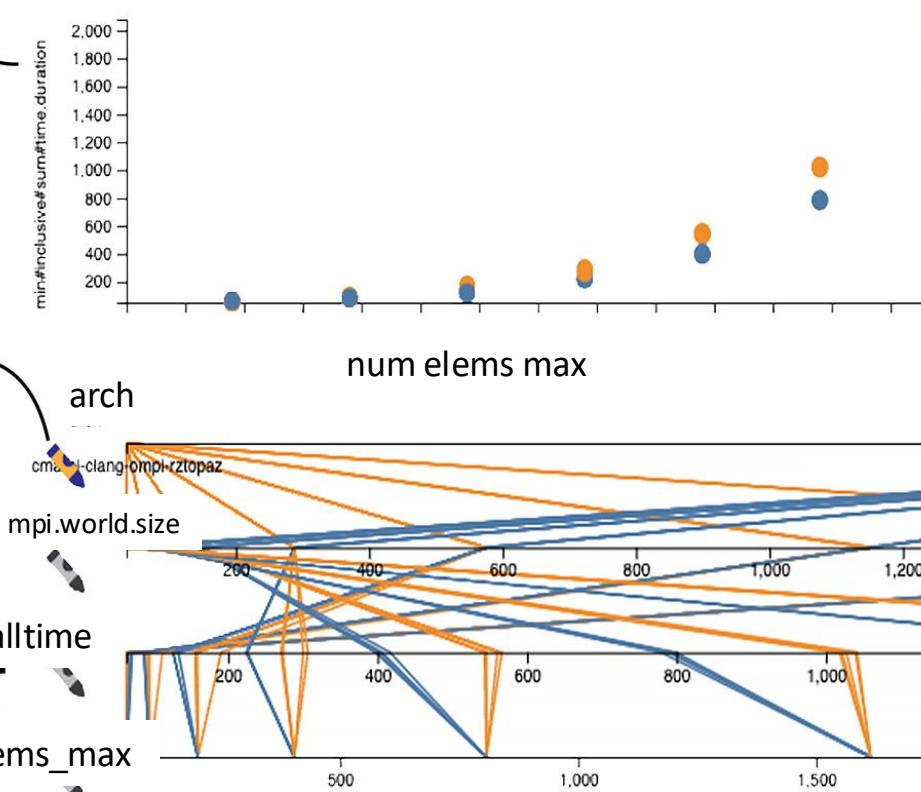
The metric values are associated with one node in the call tree.

Clicking the crayon separates data by architecture

Parallel lines show correlation between program runtime and number of simulated elements

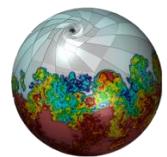
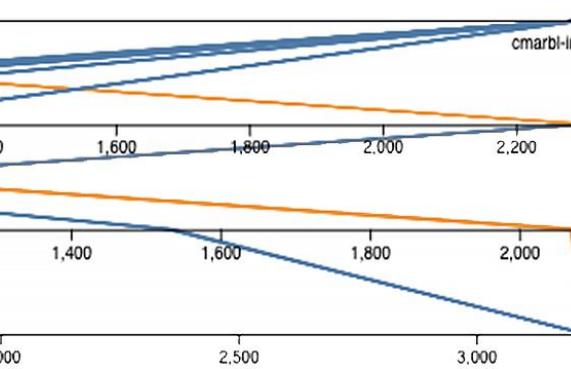
walltime

num\_elems\_max

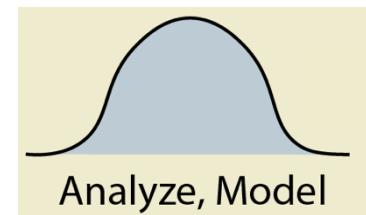


Each point represents a profile. All profiles are currently selected.

Criss-crossing lines show inverse correlation between number of MPI threads and program runtime

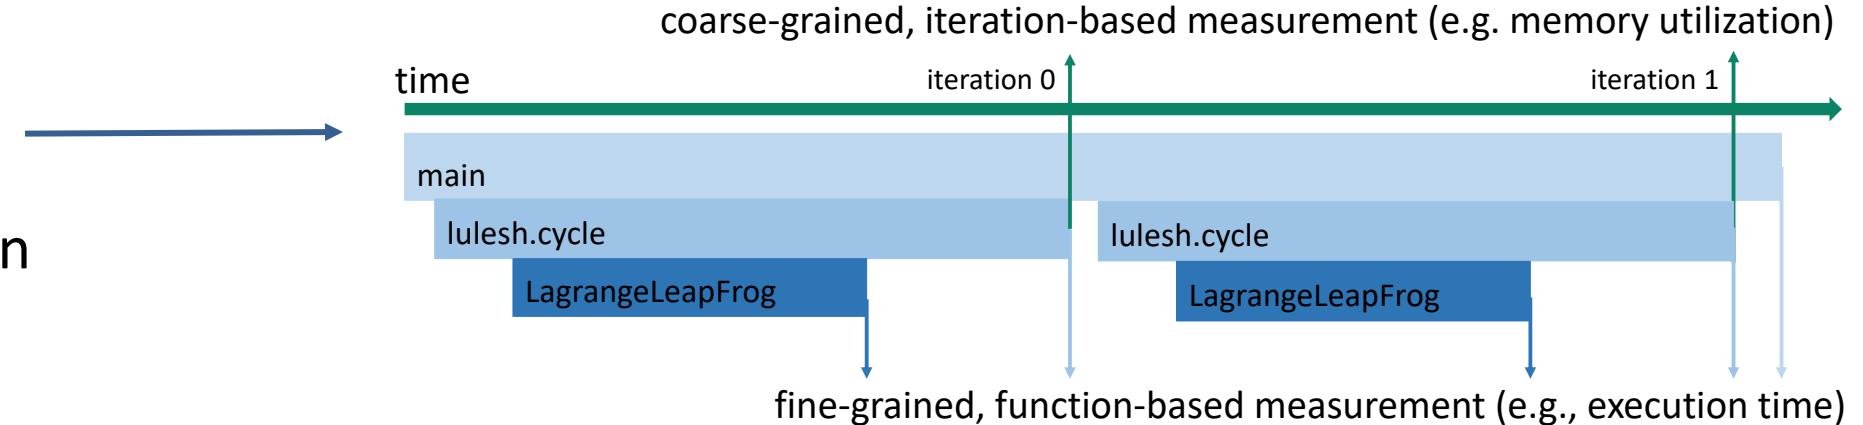


# Can we observe performance fluctuations over time?



- Caliper collects metrics at set intervals
- Thicket can then categorize temporal patterns [1]:

$$P_{temporal}(t) = 1 - \frac{\sum_{t=0}^T M_t}{\sum_{t=0}^T \max_{0 < t < T} M_t}$$



Pattern	Constant	Phased	Dynamic	Sporadic
Score	0.0-0.2	0.2-0.4	0.4-0.6	0.6-1.0
Symbol	→	~	↔	↗

[1] I.B. Peng, I. Karlin, M. B. Gokhale, K. Shoga, M. P. LeGendre, and T. Gamblin. A holistic view of memory utilization on hpc systems: Current and future trends. Proceedings of the International Symposium on Memory Systems, 2021.

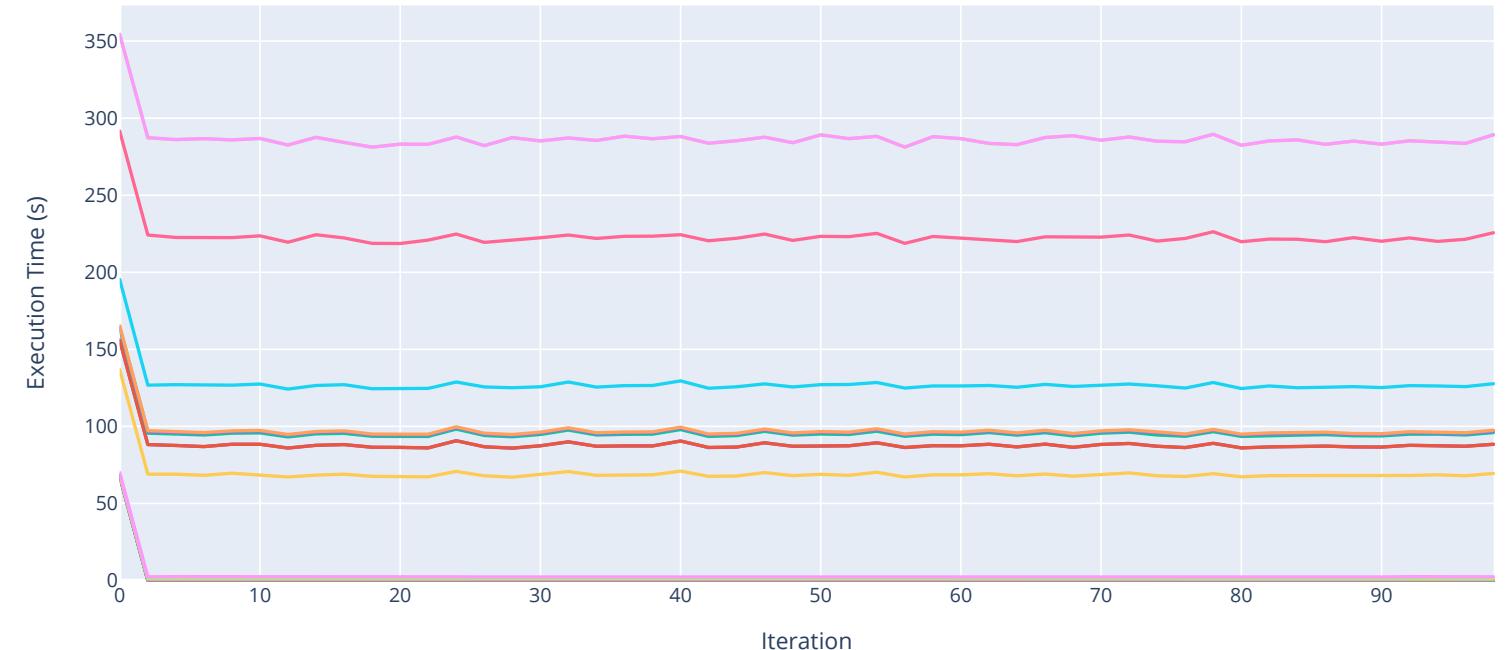


# Timeseries metrics visualizations

- Display pattern symbol and temporal score as part of the call tree
- Use python plotting libraries to create a more granular visualization

0.488 ↗ lulesh.cycle  
└ 0.280 ↗ LagrangeLeapFrog  
  └ 0.242 ↗ CalcTimeConstraintsForElems  
    └ 0.151 ↗ LagrangeElements

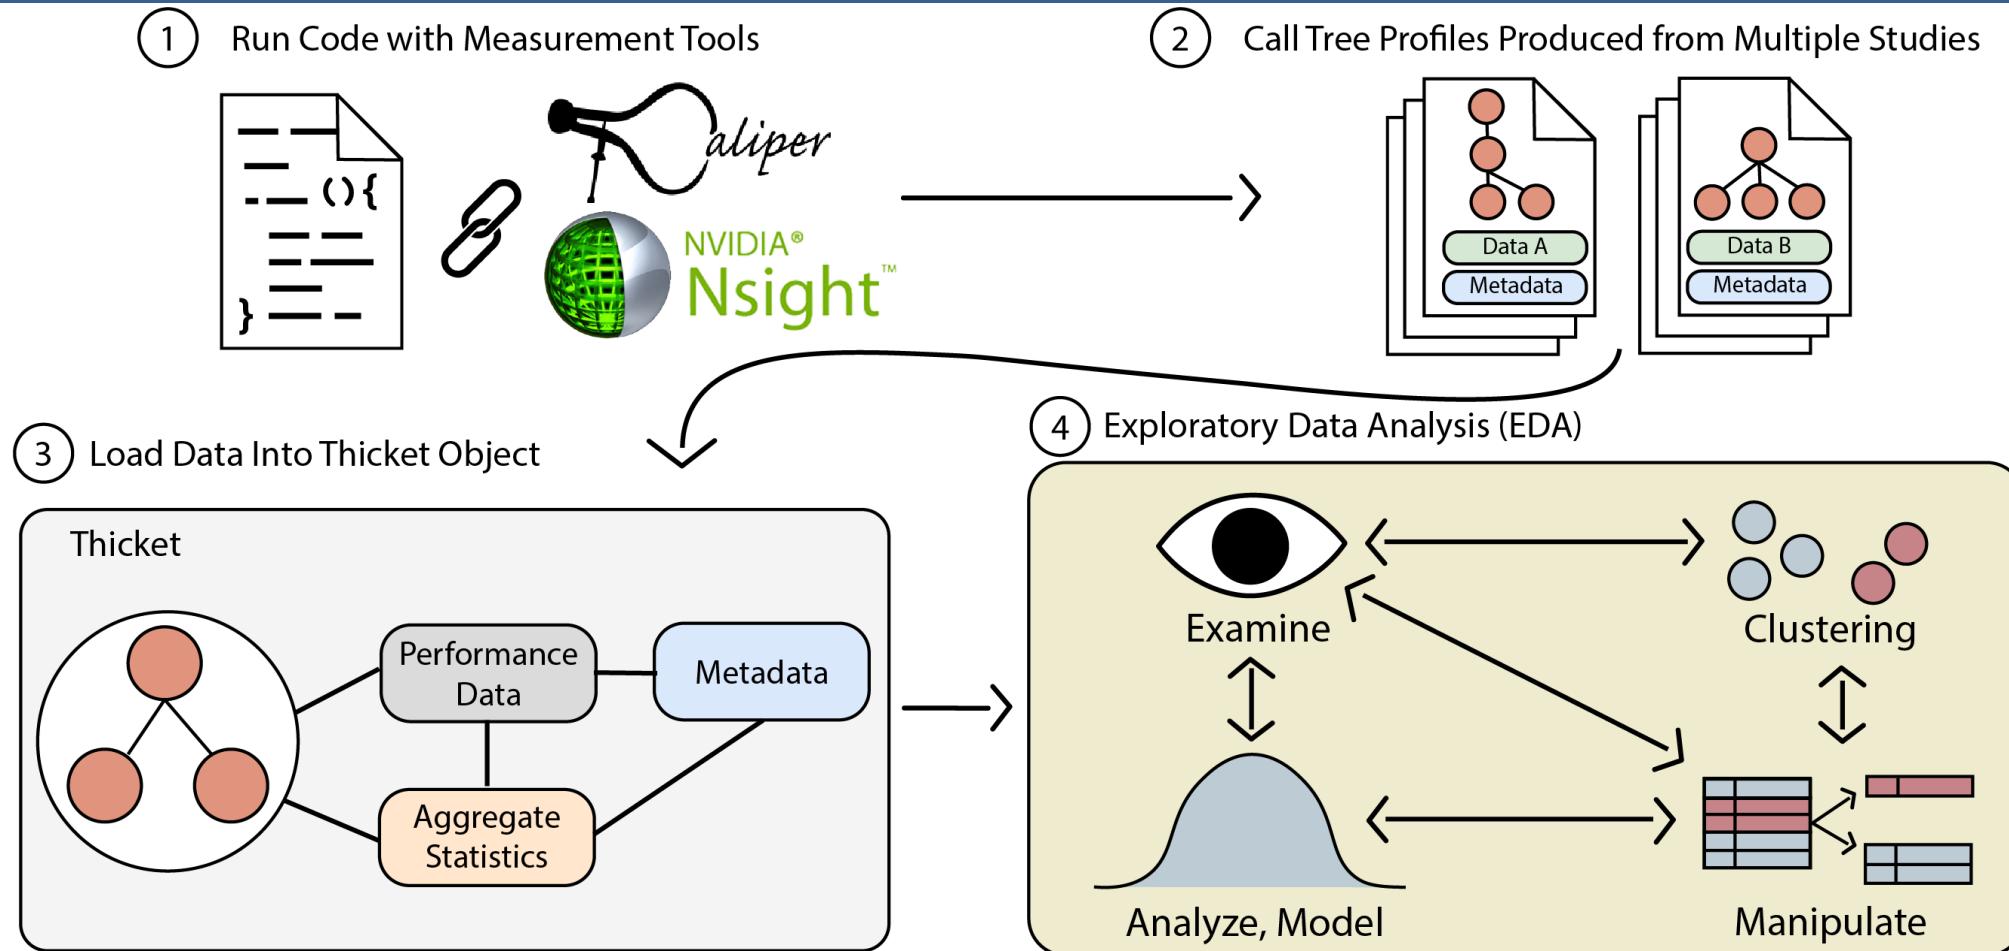
67788.343 → lulesh.cycle  
└ 8168.648 LagrangeLeapFrog  
  └ 8783.894 CalcTimeConstraintsForElem  
    └ 398433.500 LagrangeElements



Function:

main	lulesh.cycle	LagrangeLeapFrog	CalcTimeConstraintsForElems	LagrangeElements
ApplyMaterialPropertiesForElems	EvalEOSForElems	CalcEnergyForElems	CalcLagrangeElements	
CalcKinematicsForElems	CalcQForElems	CalcMonotonicQForElems	LagrangeNodal	CalcForceForNodes
CalcVolumeForceForElems	CalcHourglassControlForElems	CalcFBHourglassForceForElems	IntegrateStressForElems	
				TimeIncrement

# Thicket is a toolkit for exploratory data analysis of multi-run data



<https://github.com/LLNL/thicket>

<https://github.com/LLNL/thicket-tutorial>

## Tutorial Instances: <http://bit.ly/4kGQDlc>

---

- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
  - Pre-installed Thicket, Caliper, and Benchpark and required dependencies
  - Caliper source code demos
  - Thicket Jupyter notebooks and datasets for performance analysis



When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
  - First initial followed by last name (e.g., jdoe)
- PW: hpctutorial25



# CASC

Center for Applied  
Scientific Computing



**Lawrence Livermore  
National Laboratory**

#### **Disclaimer**

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

# Reproducible Experiments with Benchpark

Stephanie Brink



# We benchmark HPC systems for a variety of reasons

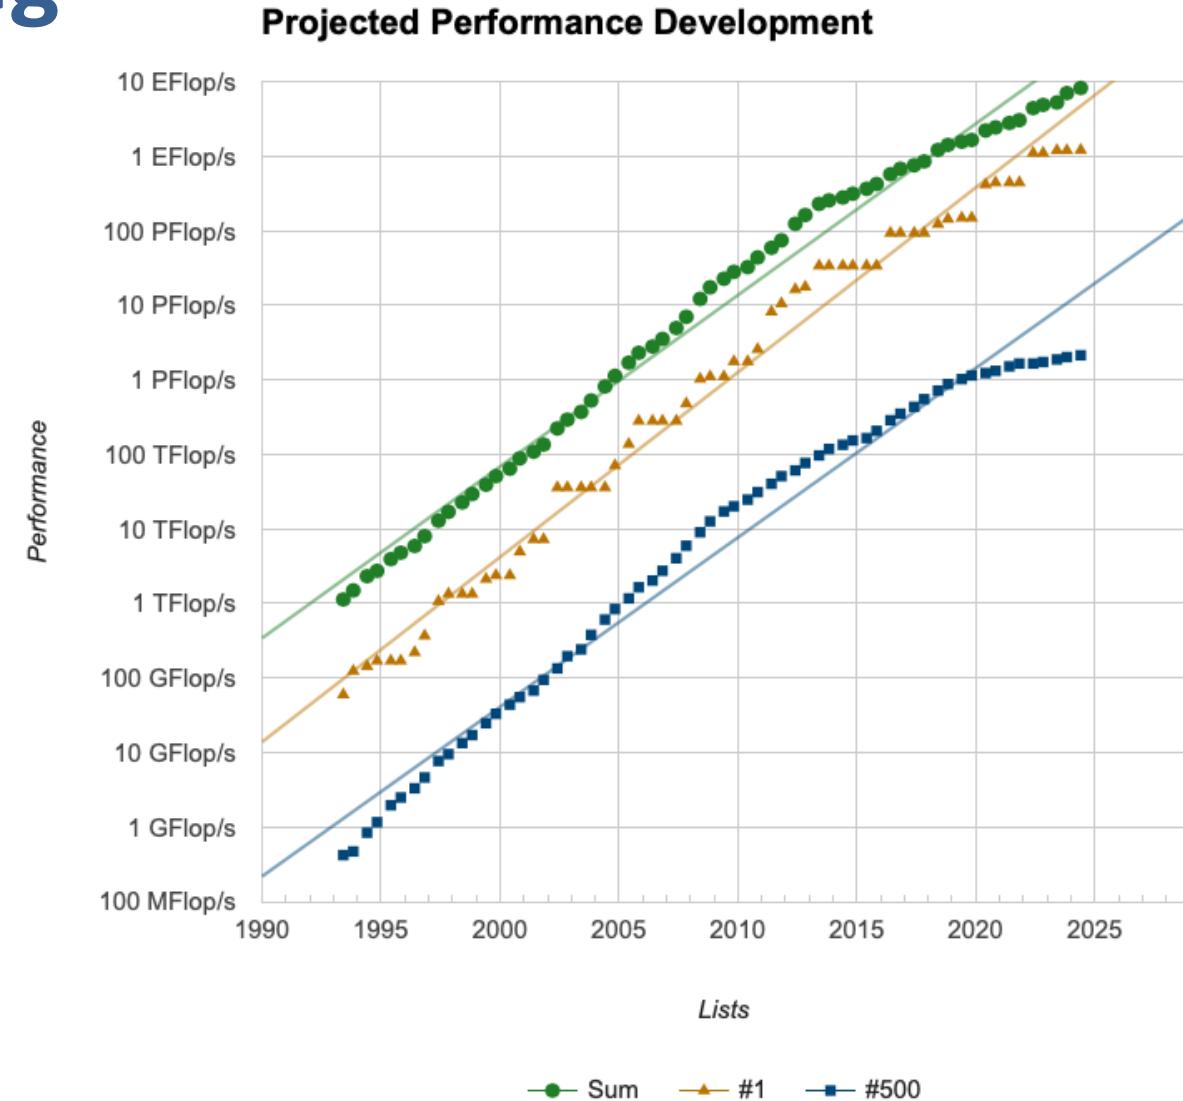
- Procurements
  - Communicate datacenter workload to vendors
  - Co-design systems, monitor progress
  - System acceptance (contractual specification)
- Validation of software stack, tools
  - Compilers
  - Debuggers
  - Correctness tools
  - Performance tools
- Research
  - Programming models
  - Computational methods
  - Performance across architectures



Many stakeholders, communication is key for collaboration

# Benchmarking is challenging

- What are we trying to characterize?
- Are we capturing the best the system can do?
- Is something else impacting performance?
- Did we build and run the code in the *optimal* and *reproducible* way?



# HPC benchmarks run on diverse HPC systems



Lawrence Berkeley Nat'l Lab  
AMD Zen + NVIDIA



Lawrence Livermore Nat'l Lab  
IBM Power9 + NVIDIA



Oak Ridge National Lab  
AMD Zen + Radeon



Argonne National Lab  
Intel Xeon + Xe

- Benchmark source code
  - Abstraction (OpenMP, RAJA, Kokkos)
  - Hardware-specific (CUDA, ROCm)
- Optimized code for the CPU and GPU
  - Must make effective use of the hardware
  - Can make 10-100x performance difference
- Rely heavily on system packages
  - Need to use optimized communication and MPI libraries that come with machines

# Writing benchmark source code is only the beginning

**State of the practice:**  
*HPC system benchmarking is manual!*



- **Building** on each system is different, porting the builds to new systems is manual
- **Running** on each system is different, porting run scripts to new systems is manual
- Systems keep **changing**, requiring updates to how we build and run benchmarks
- **Triggering** builds and runs is manual: benchmark results don't stay up to date
- **Performance analysis** of results is manual

Communicate technical specification via code, documentation, white papers



Lawrence Livermore National Laboratory  
LLNL-CFPRES-2202755

Join #benchpark-support on slack.spack.io

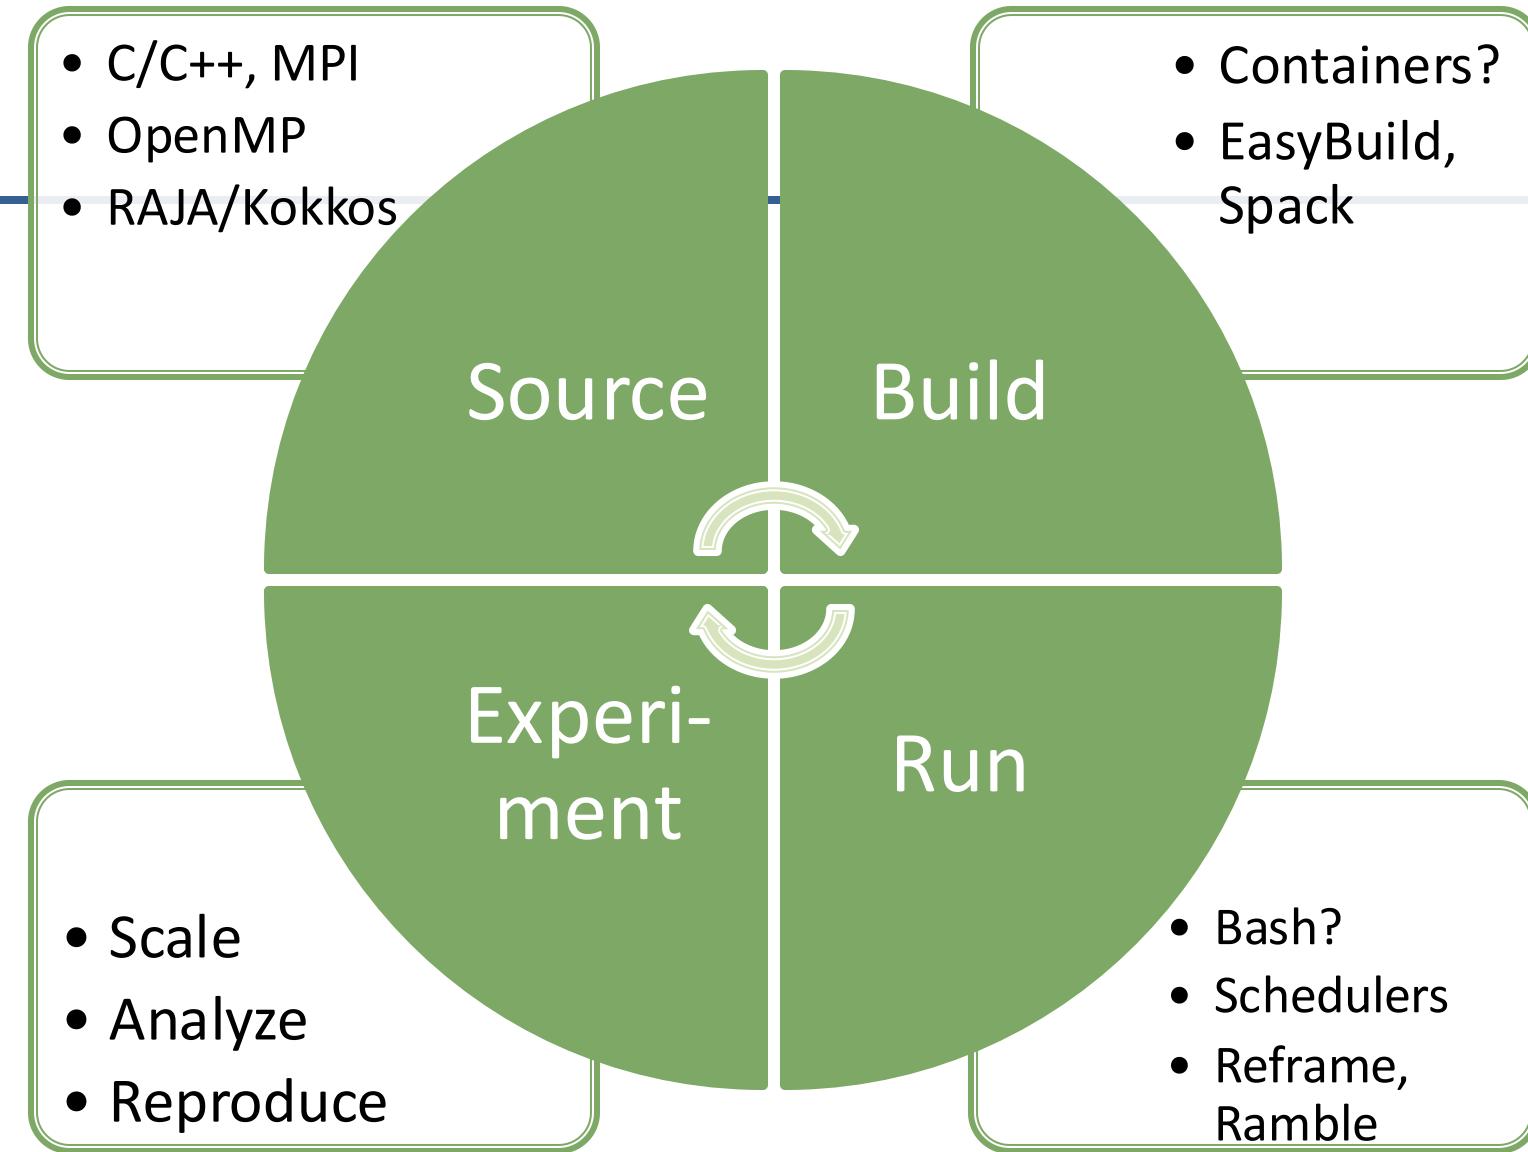


<https://github.com/llnl/benchpark>



# HPC Benchmarks are HPC Software

- Portability
- Maintenance
- Testing/CI
- Verification
- Reproducibility



All components must work *for your system*, focus on explainable performance

# Benchpark enables complete specification of HPC benchmarks



*Infrastructure-as-code* benchmark specification codifies:

- Benchmark build and run instructions
- **HPC Systems**
- **HPC Experiments**



[github.com/Llnl/benchpark](https://github.com/Llnl/benchpark)



Leverage advances in HPC automation

- Source code
- Build specification
- Run specification
- CI



Ramble  
Spack



**Every part of the specification codified: use to communicate, automate**



Lawrence Livermore National Laboratory  
LLNL-CFRES-2202755

Join #benchpark-support on slack.spack.io



<https://github.com/Llnl/benchpark>



# Benchpark enables *reproducible* specifications of benchmarks

## Specify

- How to build and run benchmarks on a system

## Run

- Run an experiment on a system

## Reproduce

- Re-run an experiment on a system

## Replicate

- Run an experiment on a new system

## Maintain

- CI: Run experiment on HPC systems

## Record

- Perf.measurements + *full spec* of experiment

Full specification enables reproducibility, replicability, and automation

# HPC System definition for Performance

- Resources (CPU cores, GPUs)
- Scheduler (Slurm, flux)
- Process mapping (cores, sockets, GPUs, NICs)
- On-node parallelism (CUDA, ROCM, OpenMP)
- Software stack
  - Compilers (which is best?)
  - MPI implementations (best?)
  - Math libraries



**Goal: Use the system correctly**

# HPC System in Benchpark: *Specify Once*

- ✓  systems
  - >  all\_hardware\_descriptions
  - >  aws-pcluster
  - >  generic-x86
  - >  lanl-venado
  - >  llnl-cluster
  - >  llnl-elcapitan
  - >  llnl-sierra
  - >  riken-fugaku

```
class LlnlElcapitan(System):  
    variant(  
        "rocm",  
        default="5.5.1",  
        values=("5.4.3", "5.5.1", "6.2.4"),  
        description="ROCM version",  
    )  
  
    variant(  
        "compiler",  
        default="cce",  
        values=("gcc", "cce", "rocmcc"),  
        description="Which compiler to use",  
    )  
  
    def initialize(self):  
        super().initialize()  
        self.scheduler = "flux"  
        self.sys_cores_per_node = "128"
```

**benchpark system init --dest=elcap llnl-elcapitan rocm=6.2.4 compiler=cce**



Search docs

**BASICS**[For the Impatient](#)[Getting Started](#)[Benchpark Commands](#)[Benchpark Workflow](#)[Frequently Asked Questions](#)**CATALOGUE**[System Specifications](#)[Benchmarks and Experiments](#)**TUTORIALS**[Hello Benchpark Example](#)[Running on an LLNL System](#)[Comparing two Experiments Within Benchpark](#)**USING BENCHPARK**[Setting Up a Benchpark Workspace](#)[Building an Experiment in Benchpark](#)[Running an Experiment in Benchpark](#)[Experiment pass/fail and FOMs](#)[Canned Analysis for Scaling Studies](#)[Benchpark Modifiers](#)

# System Specifications

The table below provides a directory of information for systems that have been specified in Benchpark. The column headers in the table below are available for use as the `system` parameter in [benchpark setup](#).

Search: 

name	integrator.vendor	integrator.name	processor.vendor	processor.name	processor.ISA	processor.w
Atos-zen2-A100-Infiniband	Atos	XH2000	AMD	EPYC-Zen2	x86_64	zen2
AWS_PCluster-zen-EFA	AWS	ParallelCluster	AMD	EPYC-Zen	x86_64	zen
DELL-cascadelake-InfiniBand	DELL		Intel	Xeon6248R	x86_64	cascadelake
DELL-sapphirerapids-OmniPath	DELLEMCR	PowerEdge	Intel	XeonPlatinum8480	x86_64	sapphirerapi
Fujitsu-A64FX-TofuD	Fujitsu	FX1000	Fujitsu	A64FX	Armv8.2-A-SVE	aarch64
HPECray-haswell-P100-Infiniband	HPECray		Intel	Xeon-E5-2650v3	x86_64	haswell
HPECray-neoverse-H100-Slingshot	HPECray	EX254n	NVIDIA	Grace	Armv9	neoverse
HPECray-zen2-Slingshot	HPECray		AMD	EPYC-7742	x86_64	zen2
HPECray-zen3-MI250X-Slingshot	HPECray	EX235a	AMD	EPYC-Zen3	x86_64	zen3
HPECray-zen4-MI300A-Slingshot	HPECray	EX255a	AMD	EPYC-Zen4	x86_64	zen4
IBM-power9-V100-Infiniband	IBM	AC922	IBM	POWER9	ppc64le	power9
Penguin-icelake-OmniPath	PenguinComputing	RelionCluster	Intel	XeonPlatinum924248C	x86_64	icelake
Supermicro-icelake-OmniPath	Supermicro		Intel	XeonPlatinum8276L	x86_64	icelake
x86_64					x86_64	

# Systems in Benchpark: July 2025

- 4 in Europe
- 6 in US labs
- 1 in Japan
- 1 at a university
- 2 cloud systems

 all\_hardware\_descriptions

 aws-pcluster

 common

 csc-lumi

 cscs-daint

 cscs-eiger

 generic-x86

 jsc-juwels

 lanl-venado

 llnl-cluster

 llnl-elcapitan

 llnl-sierra

 riken-fugaku



# HPC Experiment definition for *Performance*

- On-node parallelism  
(CUDA, ROCM, OpenMP)
- Problem sizes
  - Overall problem size, or
  - Per node or per GPU
- Scaling studies
  - How to scale
  - How to decompose
- Resources (cores, GPUs)



**Goal: Specify reproducible sets of experiments that map onto specific Systems**

# HPC Experiment in Benchpark: *Specify Once*

```
class Amg2023(Experiment, OpenMPExperiment, CudaExperiment, ROCmExperiment,
    Scaling(ScalingMode.Strong, ScalingMode.Weak, ScalingMode.Throughput,
        Caliper)):

    variant(
        "workload",
        default="problem1",
        values=("problem1", "problem2"),
        description="problem1 or problem2",
    )

    def compute_applications_section(self):
        self.register_scaling_config({
            ScalingMode.Strong: {
                "n_resources": lambda var, itr, dim, scaling_factor: var.val(dim) * scale,
                "problem_size": lambda var, itr, dim, scaling_factor: var.val(dim) // scale,
            },
            ScalingMode.Weak: {
                "n_resources": lambda var, itr, dim, scaling_factor: var.val(dim) * scale,
                "problem_size": lambda var, itr, dim, scaling_factor: var.val(dim),
            },
            ScalingMode.Throughput: {
                "n_resources": lambda var, itr, dim, scaling_factor: var.val(dim),
                "problem_size": lambda var, itr, dim, scaling_factor: var.val(dim) * scale,
            }
        })

```

amg2023 +rocm scaling=strong workload=problem2 caliper=mpi,time

**BASICS**[For the Impatient](#)[Getting Started](#)[Benchpark Commands](#)[Benchpark Workflow](#)[Frequently Asked Questions](#)**CATALOGUE**[System Specifications](#)[Benchmarks and Experiments](#)**TUTORIALS**[Hello Benchpark Example](#)[Running on an LLNL System](#)[Comparing two Experiments Within Benchpark](#)**USING BENCHPARK**[Setting Up a Benchpark Workspace](#)[Building an Experiment in Benchpark](#)

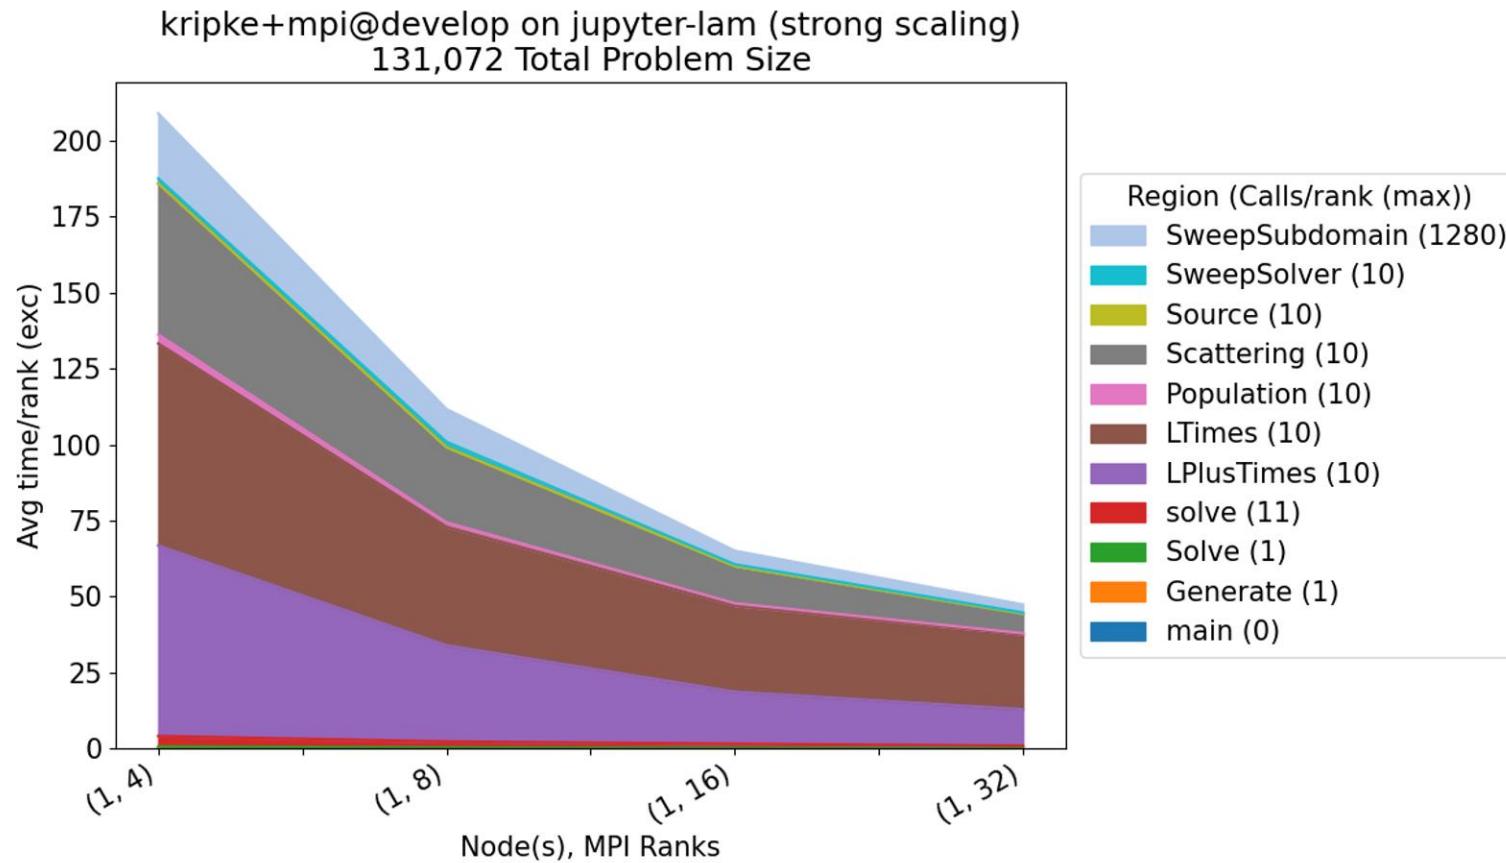
# Benchmarks and Experiments

	ad	amg2023
application-domain	[]	['engineering']
benchmark-scale	[]	['large-scale', 'multi-node', 'single-node', 'sub-node']
communication	['mpi']	['mpi']
communication-performance-characteristics	[]	['network-collectives', 'network-latency-bound']
compute-performance-characteristics	[]	['high-branching', 'mixed-precision']
math-libraries	[]	['hypre']
memory-access-characteristics	[]	['high-memory-bandwidth', 'irregular-memory-access', 'large-memory-footprint', 'regular-m']
mesh-representation	[]	['block-structured-grid']
method-type	['compiler-transformation']	['solver', 'sparse-linear-algebra']
programming-language	['c', 'c++']	['c']
programming-model	[]	['openmp', 'cuda', 'rocm']
instrumented-caliper	False	True
scaling-experiments	[]	['strong', 'weak']

# Experiments in Benchpark: July 2025

✓  experiments	>  hpcg	>  osu-micro-benchmarks	>  saxpy
>  ad	>  hpl	>  phloem	>  smb
>  amg2023	>  ior	>  quicksilver	>  stream
>  babelstream	>  kripke	>  qws	
>  genesis	>  laghos	>  raja-perf	
>  gpcnet	>  lammmps	>  remhos	
>  gromacs	>  md-test	>  salmon-tddft	
■ HPL, HPCG		■ 8 US	
■ 4 microbenchmarks		■ 1 Europe	
■ 3 MPI benchmarks		■ 2 Japan	

# `benchpark analyze` for generating pre-defined analysis charts



```
main
|- Generate
|   |- MPI_Allreduce
|   |- MPI_Comm_split
|   |- MPI_Scan
|- MPI_Allreduce
|- MPI_Bcast
|- MPI_Comm_dup
|- MPI_Comm_free
|- MPI_Comm_split
|- MPI_Finalize
|- MPI_Finalized
|- MPI_Gather
|- MPI_Get_library_version
|- MPI_Initialized
|- Solve
  |- solve
    |- LPlusTimes
    |- LTimes
    |- Population
      |- MPI_Allreduce
    |- Scattering
    |- Source
    |- SweepSolver
      |- MPI_Irecv
      |- MPI_Isend
      |- MPI_Testany
      |- MPI_Waitall
    |- SweepSubdomain
```

See <https://software.llnl.gov/benchpark/benchpark-analyze.html>

# Contributions from 11 orgs (60% non-LLNL)

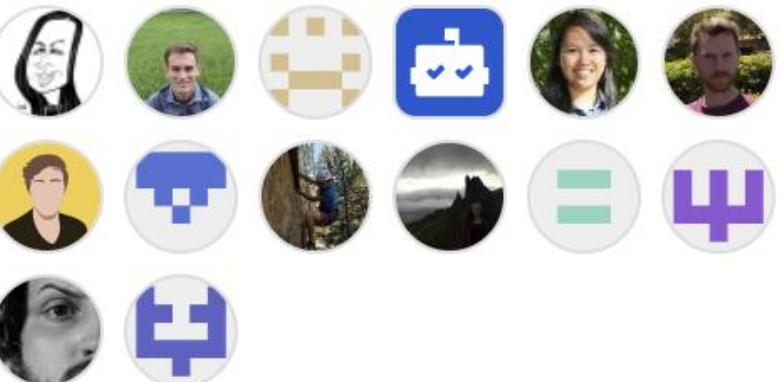
LLNL / benchpark

Type  to search

Code Issues 80 Pull requests 31 Discussions Actions Projects 3 Security 2 Insights Settings

Pulse Contributors Community

**Contributors** 29



+ 15 contributors

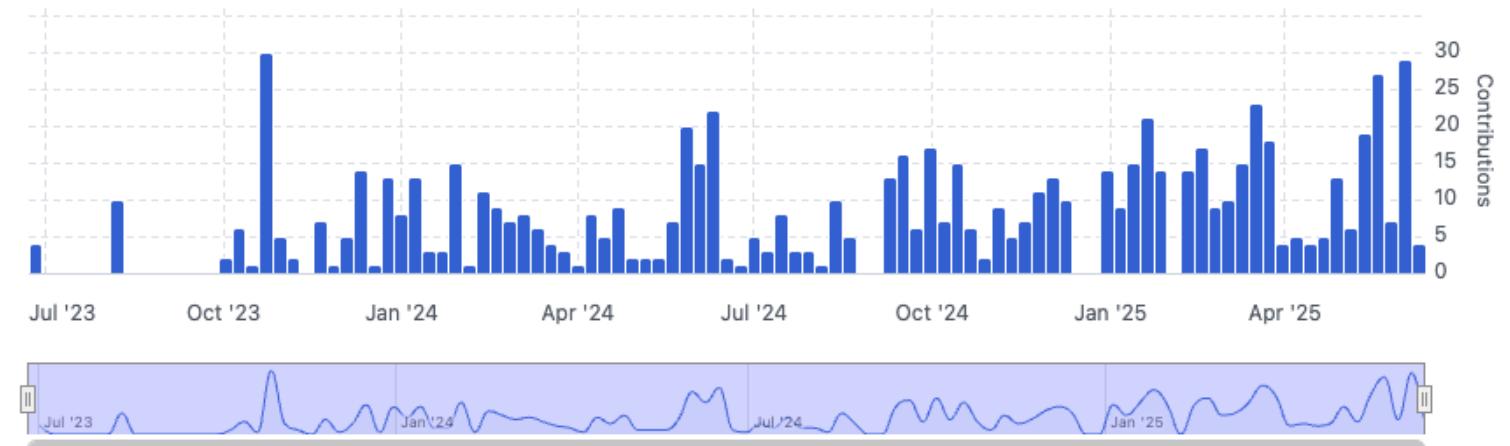
**Contributors**

Contributions per week to develop, excluding merge commits

Period: Last 24 months Contributions: Commits

**Commits over time**

Weekly from Jun 17, 2023 to Jun 14, 2025



Jul '23 Oct '23 Jan '24 Apr '24 Jul '24 Oct '24 Jan '25 Apr '25

30 25 20 15 10 5 0

LLNL-CFRES-2202755

JOIN #benchpark-support on slack.spack.io

<https://github.com/llnl/benchpark>

NNSA National Nuclear Security Administration

88

# Benchpark codifies benchmarking steps

- Benchpark does **not** replace benchmark source, build system, or Spack package
- Benchpark manages benchmark **experiments** and how they map onto **systems** with specified (or default)
  - Compilers
  - MPI/comm libraries
  - Math libraries
- Start running benchmark experiments on your system with just a few commands

✉ git clone git@github.com:LLNL/benchpark.git

🖥 benchpark list systems  
benchpark system init --dest=elcap llnl-elcapitan

👤 benchpark list experiments  
👤 benchpark experiment init dest=amg  
amg2023 +rocm scaling=strong

💻 benchpark setup amg elcap workspace  
ramble workspace setup

🐦 ramble on



# Who can use Benchpark

---

**People who want to use or distribute benchmarks for HPC!**

## **1. End Users of HPC Benchmarks**

- Install, run, analyze performance of HPC benchmarks

## **2. Benchmark Developers**

- People who want to share their benchmarks

## **3. Procurement teams at HPC Centers**

- Curate workload representation, evaluate and monitor system progress

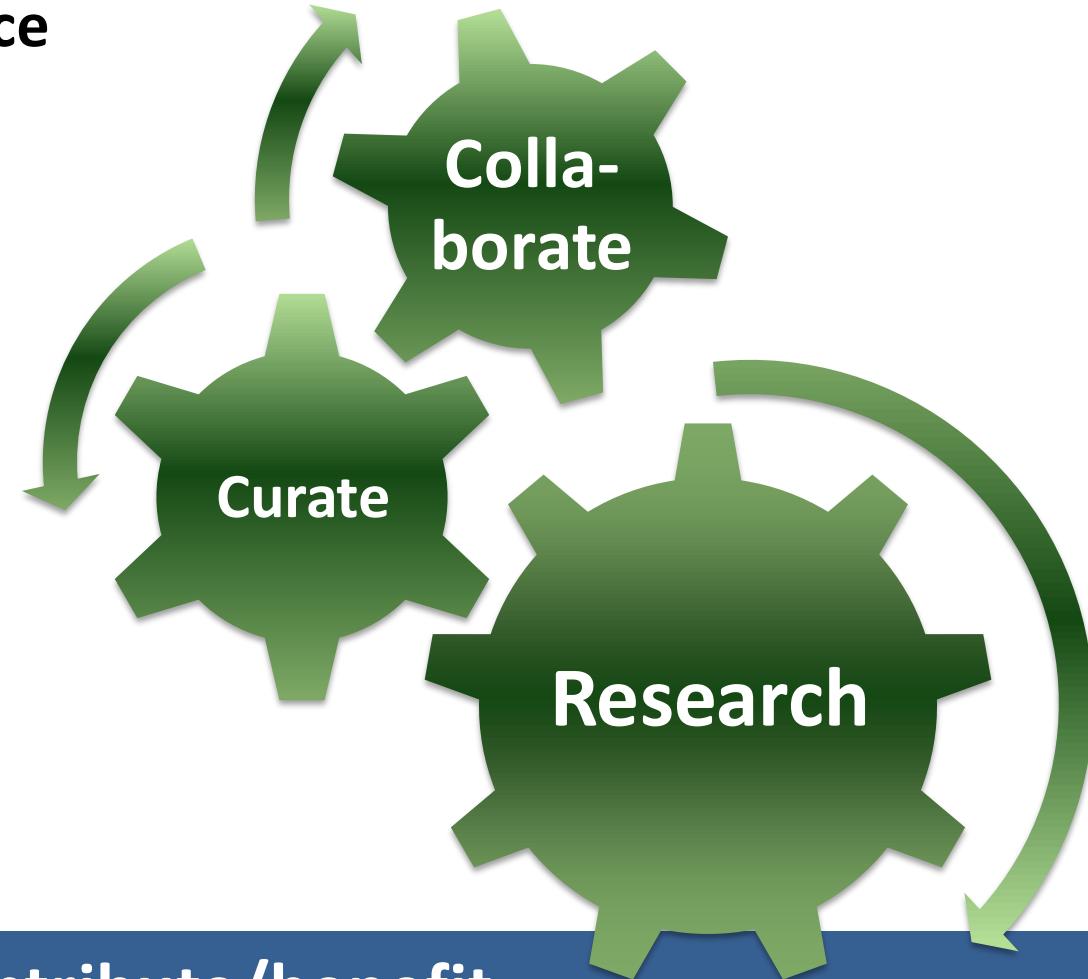
## **4. HPC Vendors**

- Understand the curated workload of HPC centers, propose systems

# Catalogued library of *working* benchmarks



- Enables exploration of large configuration space
  - Architectural
  - Software stack
  - Temporal
- What architecture and system configuration is best for *my benchmark*?
- On *my system*, is OpenMPI good enough?
- What purpose will Benchpark help you address?



Building a community to contribute/benefit

# Benchpark: Open collaborative repository for *reproducible* specifications of HPC benchmarks



*Infrastructure-as-code* benchmark specification enables **reproducibility, replicability, and automation**

- **HPC Systems**
- **HPC Experiments**

## Specify

- How to build and run benchmarks on a system

## Run

- Run an experiment on a system

## Reproduce

- Re-run an experiment on a system

## Replicate

- Run an experiment on a new system

## Maintain

- CI: Run experiment on HPC systems

## Record

- Perf.measurements + *full spec* of experiment

- Tagging, keywords for publications
- Performance metrics, metrics of usefulness
- Dashboards: Archive specs+results

- CI pipelines on PRs from GitHub at data centers across the world and in the cloud

\*Olga Pearce et al, Continuous Benchmarking, HPCTests SC|23

\*Olga Pearce et al, Repeat, Reproduce, ACM REP 2025

## Full specification enables reproducibility, replicability, and automation

# Benchpark roadmap and community engagement

---

## Future directions:

- Suite curation: Reproducible specification of an entire suite
- Tagging: Keywords for publications, finding benchmarks
- Metrics: Performance, usefulness
- Dashboards: Archive+share *specs+results*, Slices in configuration space
- CI pipelines at data centers across the world and in the cloud

## Community Engagement:

- Co-design / vendors on board, but incentive for app teams? (carrot or stick?)
- Who owns which parts of the specification and approves changes?
- Who finances R&D and maintenance?
- ROI for the people working on it? → think about post docs, researchers, etc.

**Collaboration, reproducibility, fully specified public results**



Lawrence Livermore National Laboratory  
LLNL-CFPRES-2202755

Join #benchpark-support on slack.spack.io



<https://github.com/llnl/benchpark>



# Benchpark Tutorial Materials

---

Find these slides and associated scripts here:

**[software.llnl.gov/benchpark/basic-tutorial.html](https://software.llnl.gov/benchpark/basic-tutorial.html)**

We also have a channel on Spack slack.

You can join here:

**[slack.spack.io](https://slack.spack.io)**

**Join the #benchpark-support channel!**

You can continue to ask questions here after the conference has ended.

## Tutorial Instances: <http://bit.ly/4kGQDlc>

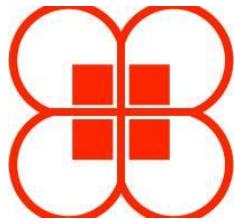
---

- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
  - Pre-installed Thicket, Caliper, and Benchpark and required dependencies
  - Caliper source code demos
  - Thicket Jupyter notebooks and datasets for performance analysis



When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
  - First initial followed by last name (e.g., jdoe)
- PW: hpctutorial25



## Join us after HPDC!

- All tutorial materials
  - [software.llnl.gov/Caliper](http://software.llnl.gov/Caliper)
  - [thicket.readthedocs.io/en/latest/tutorial\\_materials.html](https://thicket.readthedocs.io/en/latest/tutorial_materials.html)
  - [software.llnl.gov/benchpark/basic-tutorial.html](http://software.llnl.gov/benchpark/basic-tutorial.html)

- Connect with us on Spack slack

**#benchpark-support**



[slack.spack.io](https://slack.spack.io)  
**#benchpark-support**



- We want your feedback!



# Thank you!

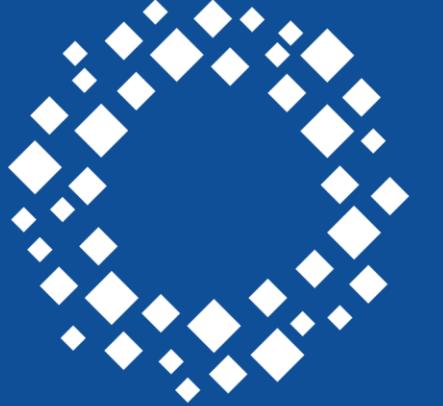
★ Star us on GitHub!

[github.com/LLNL/caliper](https://github.com/LLNL/caliper)

[github.com/LLNL/thicket](https://github.com/LLNL/thicket)

[github.com/LLNL/benchpark](https://github.com/LLNL/benchpark)





# CASC

Center for Applied  
Scientific Computing



**Lawrence Livermore  
National Laboratory**

#### **Disclaimer**

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.